# Machine Learning Script by Marko Vukotić

A Gentie introduction to Machine Learning	2
What are the main types of AI?	6
How does machine learning differ from traditional programming?	8
Overfitting and underfitting	11
Supervised vs. Unsupervised Learning	13
Regression	15
Machine Learning Terminology	16
Cross validation, Confusion Matrix, Bias and Variance	19
Essential Matrix Algebra for Neural Networks	23
Statistics in Machine Learning — Detailed Overview	31
Distributions	32
Naive Bayes	35
Probability Concepts	35
Linear Regression	38
Multiple Regression	41
Polynomial Regression	44
Normalization	45
Regularization	48
PCA (Principal Component Analysis)	50
Clustering, k-means clustering, k-nearest neighbours	53
Decision and classification trees	57
One-Hot Encoding, Label Encoding, Target Encoding, and K-Fold Cross-Validation	60
Random Forest	63
Optimizers	65
Gradient Descent	68
Stochastic gradient descent	72
Vectorization	74
AdaBoost, Gradient Boosting, and XGBoost	76
Cosine Similarity	79
Support Vector Machine	80
The Essential Main Ideas of Neural Networks	82
Train, Validate, Test	86
Activation Functions	88
Loss functions	90
Evaluation Metrics	93
The Chain Rule	97
Backpropagation Main Ideas	100
Hyperparameters	103
RELU	105
Neural Networks Multiple Inputs and Outputs	107
Neural Networks ArgMax and SoftMax	111
The SoftMax Derivative	114

Cross Entropy	117
Cross Entropy Derivatives and Backpropagation	120
Image Classification with Convolutional Neural Networks (CNNs)	123
Residuals and Residual Blocks	127
GANs	128
Tensors for Neural Networks	133
Vanishing and Exploding Gradients	136
Recurrent Neural Networks (RNNs)	137
Long Short-Term Memory (LSTM)	141
NLP	144
Bag of Words (BoW)	147
Word Embedding and Word2Vec	149
Sequence-to-Sequence (seq2seq) Encoder-Decoder Neural Networks	151
Attention	155
ELMo, BERT, Hugging Face and the others	158
Transformer Neural Networks, ChatGPT's foundation	161
Decoder-Only Transformers, ChatGPTs specific Transformer	165
Encoder-Only Transformers (like BERT) for RAG	168
Missing Data Handling	171
Learning Rate Schedulers	173
AdamW Optimizer	176
MLflow	176
Diffusion Models	177

# A Gentle Introduction to Machine Learning

# in What is Machine Learning?

Machine Learning is a subfield of artificial intelligence (AI) that enables computers to learn from data without being explicitly programmed.

Instead of writing rules, we provide examples — and the system learns the rules on its own.

# Traditional Programming vs Machine Learning

Traditional Programming	Machine Learning
You write rules manually	The system learns rules from data

Input + Rules → Output	Input + Output → Model (rules)
Good for well-defined problems	Better for complex or fuzzy patterns

# The Machine Learning Workflow

- 1. Collect Data e.g. emails labeled as spam or not spam
- Preprocess Data 
   Normalize, encode, clean
- 3. Choose Model @ e.g. decision tree, neural network, SVM
- 4. Train Model Let the model adjust internal parameters to reduce error
- 5. Evaluate Model Test accuracy, precision, recall on new data
- 6. Deploy and Improve 
  Real-world usage and periodic retraining

# Types of Machine Learning

# 1. Supervised Learning

- Learn from labeled data
- Goal: predict target/output
- Price prediction
- Image classification

# 2. Unsupervised Learning

- Learn from unlabeled data
- Goal: find hidden patterns or structures
- Clustering
- Topic modeling
- Anomaly detection

# 3. Reinforcement Learning

- Learn by interacting with an environment
- Reward for good actions, penalty for bad ones
- Game-playing AI
- Robotics
- Self-driving cars

# **Service Service Servi**

Term	Meaning
Model	A mathematical function that makes predictions

Training	Process of adjusting the model based on data
Features	Input variables used for prediction
Labels	Known output for training examples (in supervised learning)
Loss Function	Measures how wrong the model is during training
Overfitting	Model learns training data too well — bad on new data
Underfitting	Model is too simple — can't capture trends in data
Generalization	Ability to perform well on unseen data

\*\* Example: Email Spam Classifier (Supervised Learning)

1. Data: 10,000 emails labeled as "spam" or "not spam"

2. Features: Word frequencies, sender, subject line

3. Model: Logistic regression or neural network

4. Goal: Predict spam status for new emails

# Interview-Ready Takeaways

• Q: What is ML?

A: It's a way to teach computers to learn from data by optimizing mathematical models instead of writing fixed rules.

• Q: What's the difference between supervised and unsupervised learning? A: Supervised uses labeled data; unsupervised does not.

• Q: What is overfitting?

A: When a model memorizes training data and fails to generalize to new inputs.

# Summary

- ML learns patterns from data using mathematical models.
- It replaces explicit rule-writing with pattern discovery.
- V Different types (supervised, unsupervised, reinforcement) suit different problems.
- V The key challenge: generalize well to new data.

#### What are the main types of AI?

# 1. By Capability

- 1. Artificial Narrow Intelligence (ANI)
  - Definition: Specialized systems that excel at one task (e.g., image classification, language translation).
  - o Examples: Siri, AlphaGo, recommendation engines.
  - Key Point: Most Al today is narrow—cannot generalize beyond its trained scope.
- 2. Artificial General Intelligence (AGI)
  - Definition: A system with human-level cognitive abilities, able to learn and apply knowledge across domains.
  - Status: Still theoretical—no working AGI exists yet.
  - Importance: Would revolutionize automation, research, creativity.
- 3. Artificial Superintelligence (ASI)

- o Definition: An intelligence surpassing the smartest human minds in every field.
- o Speculation: Raises profound safety and ethical questions (alignment, control).

#### 2. By Functionality

#### 1. Reactive Machines

- o Trait: No memory—respond purely to current inputs.
- o Example: IBM's Deep Blue chess engine.

#### 2. Limited-Memory Systems

- o Trait: Can learn from recent data (e.g., history) to inform decisions.
- o Example: Self-driving cars use sensor data over time.

# 3. Theory of Mind (ToM) AI

- o Trait: (Hypothetical) Understands emotions, beliefs, intentions of others.
- o Goal: Enables genuine human-Al social interaction.

#### 4. Self-Aware Al

- Trait: (Speculative) Possesses self-consciousness and subjective experience.
- o Implication: Full autonomy and potentially unpredictable behavior.

#### 3. By Approach / Paradigm

<ul> <li>Symbolic (Good-Old-Fashioned Al): Logic, rules, knowledge bases (e.g., expert systems).</li> </ul>
Connectionist: Neural networks learning patterns from data.
Evolutionary: Genetic algorithms that "evolve" solutions.
Bayesian / Probabilistic: Statistical inference and uncertainty modeling.
4. Why It Matters in Interviews
Shows breadth: Understanding different axes of Al classification demonstrates deep familiarity.
Drives questions: Expect follow-ups on pros/cons, real-world examples, and safety/ethics.
5. Sample Interview Questions
• Q: "What's the difference between ANI and AGI?" A: "ANI is task-specific; AGI can generalize across tasks like a human."
<ul> <li>Q: "Give an example of a reactive vs. limited-memory AI."</li> <li>A: "Reactive: Deep Blue; Limited memory: Automotive Autopilot."</li> </ul>
<ul> <li>Q: "Why haven't we built AGI yet?"</li> <li>A: "Challenges in representation, transfer learning, common-sense reasoning, and compute constraints."</li> </ul>
How does machine learning differ from traditional programming?
1. Core Definitions

- Traditional Programming: You write explicit rules (logic + code) that map inputs to outputs.
- Machine Learning (ML): You provide data (inputs and desired outputs), and the algorithm "learns" the mapping automatically.

# 2. Process Comparison

Step	Traditional Programming	Machine Learning
1. Input	Data + Program Logic	Data (features) + Labels (ground truth)
2. Processing	Deterministic execution of code paths	Statistical model training (e.g., gradient descent)
3. Output	Predefined by your code	Model predictions based on learned patterns
4. Improvement	You manually update the code	Retrain or fine-tune on new data

# 3. Key Differences

#### Rule Creation

- o Traditional: Developer crafts every rule.
- o ML: System infers rules from examples.

# Adaptability

- o Traditional: Rigid, brittle when requirements change.
- o ML: Can adapt by retraining with updated data.

•	Comp	lexity	Hand	ling
---	------	--------	------	------

- o Traditional: Struggles with high-dimensional or unstructured data (e.g., images, speech).
- ML: Excels at uncovering patterns in large, complex datasets.

# Error Handling

- o Traditional: Errors often deterministic bugs in logic.
- o ML: Errors are statistical-models may misclassify or underfit/overfit.

#### 4. When to Use Which

- Traditional Programming:
  - o Clear, simple rules (e.g., tax calculators, string parsing).
  - o Low data volume or need for absolute correctness.
- Machine Learning:
  - o Pattern-recognition tasks (vision, NLP, recommendations).
  - o Complex, non-linear relationships in data.

#### 5. Examples

•	Traditional: If-else	logic for validating	form input (e.g.	"if age > 18.	allow registration").

- ML: Spam filter learns from thousands of labeled emails to classify new messages.
- 6. Sample Interview Questions
- Q: "Can you outline the pipeline difference between traditional and ML systems?"
   A: "Traditional: code → run → result. ML: collect/label data → train model → evaluate → deploy → monitor & retrain."
- Q: "What challenges does ML introduce that traditional programming doesn't?"
  A: "Data acquisition/labeling, model validation (overfitting vs. underfitting), feature engineering, and drift monitoring."
  - Q: "How do you decide if a task should use ML?"

A: "Assess if there's sufficient quality data, if rules are too complex to write manually, and tolerance for probabilistic errors."

#### Overfitting and underfitting

**Underfitting** 

- Happens when a model is too simple to capture the underlying pattern of the data.
- Results in poor performance on both training and test sets.
- Model has high bias and low variance.
- Example: Trying to fit a straight line to clearly nonlinear data.

# Overfitting

- Happens when a model is too complex, learning not only the pattern but also the noise in training data.
- Results in excellent training performance but poor generalization to new data.
- Model has low bias and high variance.
- Example: A very deep decision tree that memorizes training examples but fails on unseen data.

# ■ Bias-Variance Tradeoff

Aspect	Underfitting	Overfitting
Model Complexity	Too low	Too high
Training Error	High	Low
Test Error	High	High
Generalization	Poor	Poor
Bias	High	Low
Variance	Low	High

#### Phow to Handle Them

Problem	Solutions
Underfitting	Increase model complexity, train longer, add features
Overfitting	Use regularization, early stopping, more data, dropout, simpler model

# Interview Nuggets

- Q: What's underfitting?
- A: Model can't learn enough from the data; it performs poorly everywhere.
  - Q: How do you detect overfitting?

A: Training accuracy high but validation/test accuracy low.

• Q: How to fix overfitting?

A: Regularization, dropout, reduce complexity, get more data.

# Supervised vs. Unsupervised Learning

# Q € Supervised Learning

- Involves learning from labeled data each training example has an input and a corresponding correct output (label).
  - The model learns to map inputs to outputs.

- Common tasks:
  - o Classification: Assign inputs to discrete categories (e.g., dog or cat).
  - o Regression: Predict continuous values (e.g., house prices).
- Examples: Linear regression, logistic regression, SVM, neural networks.

# Unsupervised Learning

- Involves learning from unlabeled data the model tries to find patterns or structure without explicit output labels.
  - Common tasks:
    - o Clustering: Group similar data points (e.g., customer segmentation).
    - o Dimensionality reduction: Simplify data while preserving structure (e.g., PCA).
  - Examples: K-means clustering, hierarchical clustering, PCA, autoencoders.

# Key Differences

Aspect	Supervised Learning	Unsupervised Learning
Data	Labeled	Unlabeled
Goal	Predict labels or outputs	Discover structure or patterns
Feedback	Direct (from labels)	Indirect (no labels)

Touris at Alexaniah usa	Democries along the skins	
Typical Algorithms	Regression, classification	Clustering, dimensionality reduction

# Interview Nuggets

• Q: When to use supervised vs. unsupervised?

A: Use supervised when labeled data is available; use unsupervised when labels are missing or for exploratory analysis.

• Q: Can unsupervised learning be used for feature extraction? A: Yes, e.g., PCA or autoencoders to create new feature representations.

# Regression

# ₩ What is Regression?

- A supervised learning technique used to predict continuous numeric values based on input features.
- Goal: Learn a mapping from inputs X to a continuous output y with  $\epsilon$  (error), modeled as:

$$y=f(X)+\epsilon$$

# Key Types of Regression

Туре	Description	Use Cases	
Linear Regression	Models relationship as a straight line:	Predicting house prices, sales forecasting	
	y=β0+β1x1+β2x2+···+βnxn+ε		
Polynomial Regression	Extends linear regression by fitting polynomial terms for non-linear relationships	Modeling curved trends	
Ridge Regression	Linear regression with L2 regularization to prevent overfitting	Handling multicollinearity	
Lasso Regression	Linear regression with L1 regularization for feature selection	Sparse models and feature reduction	

#### How it Works

 $\bullet$  Finds parameters  $\beta$  that minimizes the difference between predicted and actual outputs, often using Mean Squared Error (MSE) as loss:

$$ext{MSE} = rac{1}{N} \sum_{i=1}^{N} \left(y_i - \hat{y}_i
ight)^2$$

• We can use methods like Ordinary Least Squares (OLS) or Gradient Descent to optimize

# Simple Python Example (Linear Regression)

from sklearn.linear\_model import LinearRegression

model = LinearRegression()

model.fit(X\_train, y\_train)

predictions = model.predict(X\_test)

# Interview Nuggets

- Q: What assumptions does linear regression make? A: Linearity, independence, homoscedasticity (constant variance), normality of errors.
- Q: How do you evaluate regression performance? A: Metrics like MSE, RMSE, MAE, R-squared.
- Q: Why use regularization like Ridge or Lasso? A: To prevent overfitting and handle multicollinearity.

# **Machine Learning Terminology**

ML Terminology Explained

#### 1. Network

- A network in AI usually refers to an interconnected system of nodes (also called neurons or units) that process and transmit information.
  - Most commonly, it means a Neural Network a computational model inspired by the human brain.

### 2. Graphs

- Visual representation of connections between nodes (e.g., neurons in neural networks or data flow in computation graphs).
  - Helps understand structure & dependencies.

#### 3. Weights

- Parameters that in models (especially neural nets) get adjusted during training.
- Control the strength of connections between neurons.

#### 4. Learning

• The process where the model improves its performance by updating parameters based on data (e.g., minimizing loss).

# 5. Model Parameters

- Internal variables (like weights and biases) learned from training data.
- Define the behavior of the model.

#### 6. Hyperparameters

Settings configured before training (e.g., learning rate, number of layers, batch size).
Not learned by the model but tuned manually or via search methods.
7. Fitting
The process of training a model on data to find the best parameters.
8. Density
In statistics/ML, refers to a probability density function describing the likelihood of variables.
Density estimation methods learn the data distribution.
9. Deep Learning
A subset of machine learning using deep neural networks with many layers to learn hierarchical features.
10. Datasets
Collections of data used for training, validating, or testing models.
11. Parametric vs. Nonparametric Models
Parametric: Have a fixed number of parameters (e.g., linear regression). Faster but limited flexibility.

Nonparametric: Number of parameters grows with data size (e.g., k-nearest neighbors). More flexible but computationally expensive.
12. Fuzzy
• Concept from fuzzy logic; deals with reasoning that is approximate rather than fixed and exact — useful for handling uncertainty in data.
13. Feature Vector
An n-dimensional vector representing input data's measurable properties/features.
14. Transfer Learning
Technique where a model trained on one task/domain is reused or fine-tuned for a different but related task.
Speeds up training and improves performance, especially with limited data.
Cross validation, Confusion Matrix, Bias and Variance
Cross-Validation
Cross-validation is a method to reliably evaluate a model's performance by splitting data multiple ways to reduce the risk of overfitting or underfitting.
Common Method: k-Fold Cross-Validation
Split the data into k equal parts (folds)
Train the model on k-1 parts, test on the remaining one

- Repeat k times, each time with a different test fold
- Average the performance metrics

Example (k=5):

Fold 1: Test, Folds 2-5: Train

Fold 2: Test, Folds 1,3-5: Train

... etc.

- ★ Why use it?
  - Better estimate of performance
  - Detect overfitting
  - Helps with hyperparameter tuning (like learning rate, regularization)
- **B** Confusion Matrix

A confusion matrix shows how well a classification model performs on labeled test data.

# For binary classification:

	Predicted Positive	Predicted Negative	
Actual Positive	✓ True Positive (TP)	X False Negative (FN)	
Actual Negative	➤ False Positive (FP)	✓ True Negative (TN)	

# Metrics Derived From It

• **Accuracy** → How many predictions were correct?:

$$\label{eq:accuracy} \text{Accuracy} = \frac{\mathit{TP} + \mathit{TN}}{\mathit{TP} + \mathit{TN} + \mathit{FP} + \mathit{FN}}$$

Precision: → How many predicted positives were actually positive?

$$\text{Precision} = \frac{\textit{TP}}{\textit{TP} + \textit{FP}}$$

• Recall (Sensitivity) → How many actual positives were captured?

$$\text{Recall} = \frac{\mathit{TP}}{\mathit{TP} + \mathit{FN}}$$

• **F1 Score** → harmonic mean of precision and recall:

$$F1 \: Score = 2 \cdot \frac{Precision \cdot Recall}{Precision + Recall}$$

Why It Matters:

- More informative than just accuracy
- Useful when classes are imbalanced (e.g., fraud detection)

#### Bias and Variance

These explain how a model learns — and why it might fail to generalize.

#### **Bias**

- Error due to wrong assumptions in the model
- $\bullet \quad \text{ High bias} \to \text{model too simple}$
- Leads to underfitting

Example: A linear model trying to fit complex data

#### Variance

- Error due to sensitivity to small fluctuations in training data
- High variance → model too complex
- Leads to overfitting

Example: A deep neural network trained on too little data

# Bias-Variance Tradeoff

Model Type	Bias	Variance	Outcome
Linear Regression	High	Low	Underfitting
Deep Neural Net	Low	High	Overfitting
Ideal Model	Low	Low	Generalization 🗸

Goal: Find the sweet spot that balances bias and variance to generalize well on unseen data.

How These Connect:

Concept	Purpose
Cross-validation	Detect overfitting / underfitting
Confusion matrix	Understand classification performance
Bias	Model too simple → underfit
Variance	Model too complex → overfit

#### **Essential Matrix Algebra for Neural Networks**

Why Matrix Algebra in Neural Networks?

Neural networks operate on tensors, and at the most basic level:

- Vectors and matrices are just 1D and 2D tensors.
- All operations (dot products, transformations, gradients) rely on linear algebra.

You can think of a neural network layer as a matrix transformation followed by a non-linear activation.

Core Matrix Concepts

1. Matrix Multiplication (Dot Product)

The heart of neural nets.

Given:

Input:

 $\mathbf{x} \in \mathbb{R}^{1 imes n} \quad ext{or} \quad (n,) \quad ext{(row vector)}$ 

Weights:

 $\mathbf{W} \in \mathbb{R}^{n imes m}$ 

Output:

 $\mathbf{y} \in \mathbb{R}^{1 imes m}$ 

Then:

 $\mathbf{y} = \mathbf{x} \mathbf{W}$ 

Each output neuron is a weighted sum of inputs.

# Example

$$A = egin{bmatrix} 1 & 2 \ 3 & 4 \end{bmatrix}, B = egin{bmatrix} 5 & 6 \ 7 & 8 \end{bmatrix}$$

Calculate  $C = A \times B$ :

• 
$$c_{11} = 1 * 5 + 2 * 7 = 5 + 14 = 19$$

• 
$$c_{12} = 1 * 6 + 2 * 8 = 6 + 16 = 22$$

• 
$$c_{21} = 3 * 5 + 4 * 7 = 15 + 28 = 43$$

• 
$$c_{22} = 3*6+4*8 = 18+32 = 50$$

Result:

$$C = egin{bmatrix} 19 & 22 \ 43 & 50 \end{bmatrix}$$

**Batch Processing** 

When inputs are stacked into a batch X of shape (batch\_size, input\_dim), then:

•	ln	ou	t b	at	ch	

$$\mathbf{X} \in \mathbb{R}^{batch\_size \times input\_dim}$$

Weights:

$$\mathbf{W} \in \mathbb{R}^{input\_dim \times output\_dim}$$

Bias:

$$\mathbf{b} \in \mathbb{R}^{1 \times output\_dim}$$

(Broadcasted across the batch dimension)

Output:

$$\mathbf{Y} \in \mathbb{R}^{batch\_size \times output\_dim}$$

# 2. Transpose

$$\left(\mathbf{A}^T
ight)_{ij} = \mathbf{A}_{ji}$$

# What this means:

- The element at row i, column j in  $\mathbf{A}^T$  is the element at row j, column i in  $\mathbf{A}$ .
- Transpose flips the matrix over its diagonal, swapping rows and columns.

# Example:

If

$$\mathbf{A} = egin{bmatrix} 1 & 2 \ 3 & 4 \end{bmatrix}$$

then

$$\mathbf{A}^T = \begin{bmatrix} 1 & 3 \\ 2 & 4 \end{bmatrix}$$

Used in:

• Backpropagation (especially when computing gradients)

- Swapping between input/output dimensions
- 3. Identity Matrix (I)

#### Definition:

A square matrix with 1's on the main diagonal and 0's elsewhere.

$$\mathbf{I} = egin{bmatrix} 1 & 0 & \cdots & 0 \ 0 & 1 & \cdots & 0 \ dots & dots & \ddots & dots \ 0 & 0 & \cdots & 1 \end{bmatrix}$$

• Acts like the number 1 for matrix multiplication:

AI=IA=A

Used in:

- Gradient computations
- Initialization tricks
- 4. Element-wise Operations

For activation functions:

- ReLU: f(x) = max(0, x)
- Sigmoid:  $\sigma(x) = 1 / (1 + e^{-x})$

Note: These are applied element-wise, not via matrix multiplication.

5. Hadamard Product (©)

Element-wise multiplication:

$$\mathbf{C} = \mathbf{A} \odot \mathbf{B}$$

Example:

If

$$\mathbf{A} = egin{bmatrix} 1 & 2 \ 3 & 4 \end{bmatrix}, \quad \mathbf{B} = egin{bmatrix} 5 & 6 \ 7 & 8 \end{bmatrix}$$

Then

$$\mathbf{C} = \mathbf{A} \odot \mathbf{B} = egin{bmatrix} 1 \cdot 5 & 2 \cdot 6 \\ 3 \cdot 7 & 4 \cdot 8 \end{bmatrix} = egin{bmatrix} 5 & 12 \\ 21 & 32 \end{bmatrix}$$

Used in:

- Gradient updates (especially for element-wise activation gradients)
- LSTM and attention models

#### 6. Matrix Inversion (A-1)

Mostly not used in forward passes (too slow and unstable), but:

- Important in theory (e.g. solving Ax = b)
- Related to normal equations in linear regression

# Neural Network Layer as Matrix Equation

For a single layer:

### Step-by-step:

1. Linear transformation:

$$z = XW + b$$

- $\mathbf{X}$ : Input data (shape: batch\_size  $\times$  input\_dim)
- W: Weights (shape: input\_dim × output\_dim)
- **b**: Bias (shape: 1 × output\_dim, broadcasted)
- 2. Activation function:

$$\mathbf{a} = f(\mathbf{z})$$

- f: Non-linear activation function (e.g., ReLU, sigmoid, tanh)
- a: Output of the layer

Combined:

$$\mathbf{a} = f(\mathbf{X}\mathbf{W} + \mathbf{b})$$

This forms the **core** of how neural networks learn non-linear functions.

#### Matrix Algebra in Backpropagation

Backprop involves computing gradient of the loss L with respect to the weights W via matrix chain rule:

$$\frac{\partial L}{\partial \mathbf{W}} = \frac{\partial L}{\partial \mathbf{a}} \cdot \frac{\partial \mathbf{a}}{\partial \mathbf{z}} \cdot \frac{\partial \mathbf{z}}{\partial \mathbf{W}}$$

Each gradient step is a matrix multiplication or Hadamard product.

The transposition of matrices is critical here (e.g., for aligning dimensions in  $X^T \cdot \delta$ ).

Example: Fully Connected Forward Pass

Given:

X: shape  $(64, 100) \rightarrow$  batch of 64 samples, 100 features each.

W (weight matrix): shape  $(100, 50) \rightarrow$  This is a **weight matrix** that transforms the 100 input features into 50 outputs (hidden neurons).

b: shape  $(50,) \rightarrow$  This is a **bias vector** for each of the 50 neurons.

Then:

 $Z = X \times W + b$  # Shape: (64, 50)

- → Matrix Multiplication: When you multiply X×W, you get a matrix Z of shape (64, 50):
  - Each of the 64 samples is multiplied by the weight matrix WWW, resulting in 50 values (one per neuron) per sample.
  - Intuitively, each of the 64 input vectors (length 100) is transformed to a 50-dimensional vector.
  - Then we add bias b.
  - b is (50,), a vector.
  - Adding b means adding the bias term to **each row** of the result.
  - This shifts the activation values for each neuron by its bias.

Result shape remains (64, 50).

And last:

A = relu(Z) # Shape: (64, 50)

- $\rightarrow$  We apply activation function:
  - ReLU (Rectified Linear Unit) is an activation function defined as:

ReLU(x)=max(0,x)

- It replaces all negative values with zero.
- This introduces **non-linearity** allowing the network to learn complex functions.

Resulting A has the same shape (64, 50).

**Example:** Backward Pass

Let:

• 
$$\delta = \frac{\partial L}{\partial \mathbf{Z}} \in \mathbb{R}^{64 \times 50}$$

- This comes from the gradient of loss through the activation function
- $ullet \ \delta = rac{\partial L}{\partial \mathbf{A}} \odot f'(Z) \leftarrow$  element-wise (Hadamard) product

# **Gradients:**

• With respect to weights:

$$\frac{\partial L}{\partial \mathbf{W}} = \mathbf{X}^{\top} \cdot \boldsymbol{\delta} \quad \text{(Shape: } 100 \times 50\text{)}$$

• With respect to bias:

$$\frac{\partial L}{\partial \mathbf{b}} = \operatorname{sum}(\delta, \operatorname{axis} = 0) \quad \text{(Shape: 50)}$$

• With respect to inputs (for passing backward to earlier layer):

$$\frac{\partial L}{\partial \mathbf{X}} = \boldsymbol{\delta} \cdot \mathbf{W}^{\top} \quad \text{(Shape: } 64 \times 100\text{)}$$

# Matrix Algebra Terms

Concept	What It Means
Matrix Multiplication	Combines input and weights

Transpose	Swaps axes for correct dimension flow
Dot Product	Summarizes similarity or projection
Hadamard Product	Element-wise multiplication
Identity Matrix	Neutral element in multiplication
Inverse	Solves systems (rare in practice)
Rank	Determines if matrix is full-featured

# Interview-Ready Insights

• Q: Why use matrix multiplication in neural networks?

A: It efficiently applies weight transformations to all inputs in a batch using vectorized operations.

• Q: What's the role of the transpose in backprop?

A: It aligns dimensions to compute gradients correctly during weight updates.

• Q: Why do we rarely invert matrices in deep learning?

A: Inversion is computationally expensive and unstable; we use gradient-based optimization instead.

#### Statistics in Machine Learning — Detailed Overview

1. Descriptive Statistics

Used to summarize and describe the main features of a dataset.

#### Mean (Average):

$$\bar{x} = \frac{1}{n} \sum_{i=1}^n x_i$$

The central value of data.

Median: The middle value when data is sorted. Useful for skewed data.

Mode: Most frequent value.

Variance ( $\sigma^2$ ) and Standard Deviation ( $\sigma$ ):  $\sigma$  - sigma Measure data spread (dispersion).

$$\sigma^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^2$$

Range, Quartiles, IQR: Measure of data spread and outliers.

#### 2. Probability Distributions

Models uncertainty and randomness in data.

Discrete distributions: Binomial, Poisson

Continuous distributions: Normal (Gaussian), Uniform, Exponential

Normal distribution is particularly important due to the Central Limit Theorem and assumption in many ML models.

#### **Distributions**

#### **Normal Distribution**

A continuous distribution that is bell-shaped and symmetric around the mean.

$$f(x)=rac{1}{\sqrt{2\pi\sigma^2}}\,e^{-rac{(x-\mu)^2}{2\sigma^2}}$$

- μ: mean (center of the curve)
- σ<sup>2</sup>: variance (controls spread)

#### Properties:

- Mean = Median = Mode = μ
- 68% of data lies within 1σ, 95% within 2σ
- Widely used due to Central Limit Theorem: sums of random variables tend to be normally distributed.

#### ML Use:

- Assumed in Linear Regression errors
- Used in Gaussian Naive Bayes
- Basis of Gaussian Mixture Models
- Underpins z-scores, confidence intervals

#### **Bernoulli Distribution**

Distribution for a binary outcome (success/failure).

$$P(X = x) = p^{x}(1 - p)^{1-x}, \quad x \in \{0, 1\}$$

• p: probability of success

#### ML Use:

- Binary classification (e.g., logistic regression outputs)
- Evaluating classifiers with binary outcomes

#### **Binomial Distribution**

Number of successes in n independent Bernoulli trials.

$$P(X=k) = \binom{n}{k} p^k (1-p)^{n-k}$$

- n: number of trials
- p: probability of success
- k: number of successes

### ML Use:

- Hypothesis testing
- Modeling count data from binary outcomes

#### **Poisson Distribution**

Number of times an event occurs in a fixed interval, with average rate λ\lambdaλ:

$$P(X=k) = \frac{\lambda^k e^{-\lambda}}{k!}$$

#### ML Use:

- Modeling rare events (e.g., # of clicks, arrivals)
- Used in Poisson regression for count data

#### **Exponential Distribution**

Definition

Models time between events in a Poisson process:

$$f(x) = \lambda e^{-\lambda x}, \quad x \ge 0$$

NL Use:

- Time until next event (survival analysis, failure times)
- Memoryless: P(X>s+t | X>s)=P(X>t)P(X>s+t | X>s) = P(X>t)P(X>s+t | X>s)=P(X>t)

# **► Uniform Distribution**

**№** Definition

Equal probability across an interval:

$$f(x)=rac{1}{b-a},\quad a\leq x\leq b$$

NL Use:

- Weight initialization in neural nets
- Random sampling, baseline distributions

# **Distribution Table**

Distribution	Туре	ML Use-Case	Key Params
Normal	Continuous	Regression, GMM, assumptions, z-scores	μ,σ\mu, \sigmaμ,σ
Bernoulli	Discrete	Binary classification, NB, logistic regression	ppp
Binomial	Discrete	Count binary successes	n,pn, pn,p
Poisson	Discrete	Count events over time	λ\lambdaλ
Exponential	Continuous	Time until next event	λ\lambdaλ
Uniform	Continuous	Random sampling, init weights	a,ba, ba,b
Beta	Continuous	Bayesian inference on probabilities	α,β\alpha, \betaα,β
Gamma	Continuous	Bayesian priors, durations	α,β\alpha, \betaα,β
Categorical	Discrete	Multi-class classification, NLP	pip_ipi
Multinomial	Discrete	NLP, text classification	n,pin, p_in,pi

#### **Naive Bayes**

Naive Bayes is a simple probabilistic classification algorithm based on Bayes' theorem and the assumption that all features are independent of each other given the class variable. This "naive" assumption simplifies calculations, making it computationally efficient and often effective, especially with smaller datasets.

Here's a more detailed explanation:

**Key Concepts:** 

#### **Supervised Learning:**

Naive Bayes is a supervised learning algorithm, meaning it is trained on labeled data to learn the relationships between features and class labels.

#### Probabilistic Classification:

It classifies data points by calculating the probability of each class given the observed features.

#### Bayes' Theorem:

It uses Bayes' theorem to calculate the posterior probability of a class given the features.

#### **Conditional Independence:**

The core assumption is that the features are conditionally independent of each other, meaning the presence of one feature does not affect the probability of other features, given the class.

# **Computational Efficiency:**

The independence assumption makes Naive Bayes computationally fast and simple to implement.

How it Works:

#### 1. Training:

Naive Bayes is trained on a dataset where each data point has labeled features and a class label.

#### 2. Probability Calculation:

It calculates the probability of each class given the observed features using Bayes' theorem.

#### 3. Prediction:

The data point is assigned to the class with the highest posterior probability.

# Advantages:

- Simplicity and Speed: Easy to implement and computationally fast due to the independence assumption.
- Effective with Small Datasets: Can perform well even with limited training data.
- Good for Text Classification: Often used for tasks like spam detection, sentiment analysis, and text categorization.

#### **Probability Concepts**

1.Probability of an event

P(A) = Number of favorable outcomes Total number of outcomes P(A) = Total number of outcomes Number of favorable outcomes Values range from 0 to 1 P(A) = 1

P(A)=1: certain event

P(A)=0: impossible event

#### 1 2. Complement Rule

P(Ac) = 1 - P(A) P(Ac) = 1 - P(A) A c A c is the event "not A"

- + 3. Addition Rules For mutually exclusive events:  $P(A \cup B) = P(A) + P(B) P(A \cup B) = P(A) + P(B)$  For general case:  $P(A \cup B) = P(A) + P(B) P(A \cap B) P(A \cup B) = P(A) + P(B) P(A \cap B)$
- 4. Multiplication Rules For independent events:  $P(A \cap B) = P(A) \cdot P(B) P(A \cap B) = P(A) \cdot P(B)$  For dependent events:  $P(A \cap B) = P(A) \cdot P(B \mid A) P(A \cap B) = P(A) \cdot P(B \mid A) P(B \mid A) P(B \mid A)$ : probability of B given A happened
- $\geq$  5. Conditional Probability  $P(A \mid B) = P(A \cap B) P(B)$ ,  $P(B) \neq 0$   $P(A \mid B) = P(B)$   $P(A \cap B)$ , P(B) = 0 Important for reasoning under uncertainty (used in Bayesian inference)
- $\bigotimes$  6. Bayes' Theorem  $P(A \mid B) = P(B \mid A) \cdot P(A) P(B) P(A \mid B) = P(B) P(B \mid A) \cdot P(A)$  Used in: Naive Bayes classifiers, probabilistic reasoning
- 7. Total Probability Theorem If  $A \ 1$ ,  $A \ 2$ , ...,  $A \ n \ A \ 1$ ,  $A \ 2$ ,...,  $A \ n$  are mutually exclusive and exhaustive events:  $P(B) = \sum_{i=1}^{n} P(B \mid A \ i) \cdot P(A \ i) P(B) = i = 1 \sum_{i=1}^{n} P(B \mid A \ i) \cdot P(A \ i)$
- 8. Expected Value (Mean) Discrete:  $E[X] = \sum x i P(x i) E[X] = \sum x i P(x i)$  Continuous:  $E[X] = \int -\infty \infty x \cdot f(x) dx E[X] = \int -\infty \infty x \cdot f(x) dx E[X] = \int -\infty \infty x \cdot f(x) dx E[X] = \int -\infty \infty x \cdot f(x) dx E[X] = \int -\infty \infty x \cdot f(x) dx E[X] = \int -\infty \infty x \cdot f(x) dx E[X] = \int -\infty \infty x \cdot f(x) dx E[X] = \int -\infty \infty x \cdot f(x) dx E[X] = \int -\infty \infty x \cdot f(x) dx E[X] = \int -\infty \infty x \cdot f(x) dx E[X] = \int -\infty \infty x \cdot f(x) dx E[X] = \int -\infty \infty x \cdot f(x) dx E[X] = \int -\infty \infty x \cdot f(x) dx E[X] = \int -\infty \infty x \cdot f(x) dx E[X] = \int -\infty \infty x \cdot f(x) dx E[X] = \int -\infty \infty x \cdot f(x) dx E[X] = \int -\infty \infty x \cdot f(x) dx E[X] = \int -\infty \infty x \cdot f(x) dx E[X] = \int -\infty \infty x \cdot f(x) dx E[X] = \int -\infty x \cdot f(x)$
- $\S$  9. Variance and Standard Deviation Var  $(X) = E[(X \mu) 2] = E[X 2] (E[X]) 2 Var(X) = E[(X \mu) 2] = E[X 2] (E[X]) 2 SD <math>(X) = Var(X) SD(X) = Var(X)$  Measures spread of the distribution
- $\bigstar$  11. Joint and Marginal Probability Joint: probability of two things happening together P(A, B) P(A,B) Marginal: probability of a single event (sum/integrate joint probs)  $P(A) = \sum B P(A, B)$  P(A)= B  $\sum P(A,B)$
- @ 12. Entropy (Information Theory)  $H(X) = -\sum p(xi) \log 2 p(xi) H(X) = -\sum p(xi) \log 2 p(xi)$  Used to measure uncertainty or "information" Used in decision trees, feature selection, compression
- ⊗ 13. KL Divergence (Relative Entropy)  $D KL(P | | Q) = \sum P(x) \log (P(x)Q(x))$  D KL (P|Q)= $\sum P(x) \log(Q(x))$  Measures difference between 2 distributions Used in regularization, variational inference

#### 3. Bayesian Statistics

Uses Bayes' theorem to update probabilities based on evidence:

$$P(A|B) = rac{P(B|A) imes P(A)}{P(B)}$$

Foundation for Bayesian ML models and probabilistic reasoning.

#### 4. Hypothesis Testing

Used to test assumptions about data.

Null Hypothesis (H0): Default assumption.

Alternative Hypothesis (H1): Opposite claim.

- p-value: Probability of observing data given H0 is true.
- Significance level (α): Threshold to reject H0 (commonly 0.05).

Tests like t-test, chi-square test, ANOVA are used to check if differences are statistically significant.

#### 5. Confidence Intervals

Range of values to estimate a population parameter with a certain confidence (e.g., 95%).

#### 6. Correlation and Covariance

Covariance measures how two variables change together.

Correlation standardizes a covariance between -1 and 1.

This helps understand relationships between features.

#### 7. Regression Analysis

Modeling relationship between dependent and independent variables.

- Linear regression: Assumes linear relationship.
- Coefficient interpretation: Sign, magnitude indicate influence.
- Residuals: Differences between observed and predicted.

#### 8. Bias-Variance Tradeoff

Bias: Error from erroneous assumptions (underfitting).

Variance: Error from sensitivity to small fluctuations (overfitting).

Goal: balance both for best generalization.

#### 9. Sampling Methods

Random sampling, stratified sampling and bootstrapping help create representative datasets and estimate model stability.

#### Summary Table

Concept	Purpose/Use	
Mean, Median, Mode	Describe central tendency	
Variance, Std Dev	Measure data spread	
Probability	Model uncertainty	

Bayesian Stats	Update beliefs with evidence
Hypothesis Testing	Validate assumptions
Confidence Intervals	Estimate parameter range
Correlation	Measure relationship between vars
Regression	Model dependency
Classification Metrics	Evaluate classifier performance
Bias-Variance Tradeoff	Optimize model generalization
Sampling	Create representative datasets

#### **Linear Regression**

What is Linear Regression?

Linear Regression is a supervised learning algorithm used to model the relationship between input features and a continuous output.

It assumes the relationship is linear — that is, the output is a weighted sum of the inputs plus a bias (intercept).

Equation of a Line

For simple linear regression (one feature):

$$y = w \cdot x + by = w \cdot x + b$$

- x = input (feature)
- w = weight (slope)
- b = bias (intercept)
- y = predicted output

For multiple features (multivariate linear regression):

$$y = w1 \cdot x1 + w2 \cdot x2 + ... + wn \cdot xn + by = w_1 \cdot x_1 + w_2 \cdot x_2 + ... + w \\ \mathbb{N} \cdot x \\ \mathbb{N} + b \ y = w1 \cdot x1 + w2 \cdot x2 + ... + wn \cdot xn + b$$

Or more compactly with vectors:

$$y = X \cdot W + b$$

Where:

- X = input vector (or matrix for multiple samples)
- W = weight vector
- b = scalar or bias vector (broadcasted)

# **@** Goal of Training

Minimize the error between predictions and actual targets.

Common Loss Function: Mean Squared Error (MSE):

$$ext{MSE} = rac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

- y<sub>i</sub> = true output
- $\hat{y}_i$  = predicted output
- n = number of data points

# How It Learns

Two main ways to find the best weights W:

1. Analytical Solution (Normal Equation)

Solves for weights in one step (no training loop):

$$\mathbf{W} = (\mathbf{X}^{\top}\mathbf{X})^{-1}\mathbf{X}^{\top}\mathbf{y}$$

Only works well if:

- X<sup>T</sup>X is invertible
- Data is not too large (computationally expensive)

#### 2. Gradient Descent

The gradient is a vector of partial derivatives of a multivariable function.

Iteratively update weights to reduce loss:

$$\mathbf{W} = \mathbf{W} - \alpha \frac{\partial L}{\partial \mathbf{W}}$$

- $\alpha$  = learning rate
- $\partial L/\partial W$  = gradient of the loss

# Assumptions of Linear Regression

- 1. Linearity: Relationship between inputs and output is linear
- 2. Homoscedasticity: Constant variance of residuals
- 3. No multicollinearity: Features shouldn't be highly correlated
- 4. Normality: Errors are normally distributed

# **TITIE** Evaluation Metrics

R<sup>2</sup> Score (Coefficient of Determination): Measures how much variance in the output is explained by the inputs.

$$R^2 = 1 - rac{SS_{
m res}}{SS_{
m tot}}$$

- $R^2 = 1 \rightarrow perfect fit$
- $R^2 = 0 \rightarrow \text{model predicts mean}$

# Example

from sklearn.linear\_model import LinearRegression

model = LinearRegression()

model.fit(X\_train, y\_train)

predictions = model.predict(X\_test)

# Interview Nuggets

- Q: What's the difference between linear and logistic regression? A: Linear is for continuous outputs, logistic is for binary classification.
- Q: Why use gradient descent if there's a closed-form solution? A: Gradient descent works better for large datasets or when X<sup>T</sup>X is not invertible.
  - Q: What are the limitations of linear regression?

A: It assumes linear relationships and can't handle complex patterns unless transformed.

# **Multiple Regression**

Multiple Regression (Multivariate Linear Regression)

Multiple Regression is a supervised learning method that models the relationship between multiple input variables (features) and a single continuous output by fitting a linear equation.

# Mathematical Model

The predicted output y is a linear combination of multiple inputs:

$$y = w_1x_1 + w_2x_2 + \cdots + w_nx_n + b$$

Or in vector/matrix form:

$$y = \mathbf{X}\mathbf{W} + b$$

- X: input feature vector (or matrix for many samples)
- W: weights (coefficients) for each feature
- b: bias or intercept term
- y: predicted continuous value

# Goal of Multiple Regression

- Find the weight vector and bias b that minimize the prediction error over all training samples.
- The usual loss function is the Mean Squared Error (MSE):

#### How Does It Learn?

1. Normal Equation (Analytical solution):

- Computes the best-fitting weights in one go
- Computationally expensive for very large feature sets or datasets
- 2. Gradient Descent (Iterative solution):
  - $\alpha$  = learning rate
  - Update weights step-by-step to reduce error
- Assumptions (Same as Simple Linear Regression)
  - Linearity: Relationship between inputs and output is linear
  - Independence: Observations are independent
  - Homoscedasticity: Constant variance of residuals/errors
  - No multicollinearity: Features should not be highly correlated
  - Normality: Errors are normally distributed (for inference)

- Evaluation Metrics
  - R-squared (R2): Proportion of variance explained by the model
  - Adjusted R<sup>2</sup>: Adjusts for number of predictors to avoid overfitting
  - Root Mean Squared Error (RMSE): Square root of MSE, in same units as output

# Interview Nuggets

• Q: What's multicollinearity and why is it a problem?

A: When features are highly correlated, it makes it hard to isolate individual effects on output, leading to unstable estimates.

• Q: How do you handle categorical variables in multiple regression?

A: Use one-hot encoding to convert categories into numeric dummy variables.

• Q: When is multiple regression inappropriate?

A: When relationships are nonlinear or features interact in complex ways not captured by a linear model.

#### **Polynomial Regression**

Polynomial Regression is an extension of **linear regression** where the relationship between the independent variable X and the dependent variable y is modeled as an **n-th degree polynomial**.

Why use Polynomial Regression?

When the relationship between X and y is non-linear, a straight line (linear regression) won't fit well. Polynomial regression fits curves by adding polynomial terms (like X^2, X^3, ...).

For a single feature X, a polynomial regression of degree ddd looks like:

$$y = \beta_0 + \beta_1 X + \beta_2 X^2 + \beta_3 X^3 + \dots + \beta_d X^d + \varepsilon$$

Where:

- $\beta 0$ ,  $\beta 1$ , ...,  $\beta d$  are coefficients to learn.
- $\epsilon$  is the error term.

#### How does it work?

Transform the input feature X into polynomial features: X,  $X^2$ ,  $X^3$ , ...,  $X^d$ .

- Use linear regression on these transformed features.
- Even though the model is linear in coefficients  $\beta$ , it models a nonlinear relationship in X.

#### Example in Python

```
# imports from sklearn
# Sample data
X = np.array([1, 2, 3, 4, 5]).reshape(-1, 1)
y = np.array([1, 4, 9, 16, 25]) # y = x^2
# Create polynomial features (degree 2)
poly = PolynomialFeatures(degree=2)
X_poly = poly.fit_transform(X)
# Fit linear regression on polynomial features
model = LinearRegression()
model.fit(X_poly, y)
# Predict
y_pred = model.predict(X_poly)
# Plot
plt.scatter(X, y, color='blue')
plt.plot(X, y_pred, color='red')
plt.title("Polynomial Regression Degree 2")
plt.show()
```

#### Normalization

Normalization refers to the process of scaling input data so that it fits within a specific range or distribution. It helps models learn faster, converge more reliably, and perform better, especially in gradient-based learning (like neural networks).

Common techniques:

#### **Min-Max Normalization**

$$x' = rac{x - x_{\min}}{x_{\max} - x_{\min}}$$

Scales features to [0, 1]

Sensitive to outliers

#### Z-score Normalization (Standardization)

$$x' = \frac{x - \mu}{\sigma}$$

Mean = 0, Std = 1

Used when data follows (or is assumed to follow) a normal distribution

Ideal for linear models, SVMs, and neural networks.

Internal normalization techniques:

#### Batch Normalization (BatchNorm)

$$\mathrm{BN}(x) = \gamma \cdot rac{x - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}} + eta$$

Applied per batch, per feature channel

Stabilizes learning and allows higher learning rates

Often used after conv or fully connected layers

#### **Layer Normalization**

Normalizes across features per sample

Often used in transformers and RNNs

#### **Instance Normalization**

Like BatchNorm, but normalizes each **individual sample** across spatial dimensions

Popular in style transfer

#### **Group Normalization**

Splits channels into groups and normalizes within each group

Works better for small batch sizes

Why normalize?

Without Normalization	With Normalization
Uneven feature scales	Features treated equally
Slower convergence	Faster, stable training
Exploding gradients	Controlled gradients
Hard to tune learning	More robust optimization

# In Practice:

- Always normalize your data before training (especially in ML or CNNs).
- Use **Z-score normalization** unless min-max is required (e.g., image pixel inputs).
- Use **BatchNorm or LayerNorm** for deeper networks to help training converge.

Туре	Where?	Goal	Formula / Idea
Min-Max	Preprocessing	Scale to [0, 1]	x-xminxmax-xminx - x_{\min}}{x_{\max} - x_{\min}}
Z-score (Standard)	Preprocessing	Center + scale	x-μσ\frac{x - \mu}{\sigma}
BatchNorm	Inside NN	Normalize per batch	Normalize + learnable scale/shift
LayerNorm	Inside NN	Normalize per sample	Across all features
InstanceNorm	Inside CNNs	Normalize per image	Common in style transfer
GroupNorm	Inside CNNs	Normalize per channel group	Useful for small batches

# Regularization

What is Regularization?

Regularization is a technique used to reduce overfitting by adding a penalty term to the loss function during model training. It discourages the model from fitting noise or overly complex patterns in the training data, improving generalization to new data.

Why Regularize?

- Without regularization, complex models (like deep neural nets or multiple regression with many features) can memorize training data, performing poorly on unseen data.
- Regularization helps the model keep weights small or simple, preventing it from relying too much on any single feature or noise.
- Common Types of Regularization
- 1. L2 Regularization (Ridge Regression)
  - Adds the sum of squared weights to the loss function:

$$\operatorname{Loss} = \operatorname{Loss}_{original} + \lambda \sum_{j=1}^n w_j^2$$

- ullet  $\lambda$  (lambda) is the regularization parameter controlling the strength of the penalty.
- Encourages weights to be small but not zero.
- Helps with multicollinearity in regression.
- 2. L1 Regularization (Lasso Regression)
  - Adds the sum of absolute weights to the loss:

$$ext{Loss} = ext{Loss}_{original} + \lambda \sum_{j=1}^{n} |w_j|$$

- Encourages sparsity many weights become exactly zero.
- Useful for feature selection by eliminating less important features.

#### 3. Elastic Net

• Combines L1 and L2 penalties:

$$ext{Loss} = ext{Loss}_{original} + \lambda_1 \sum_{j=1}^n |w_j| + \lambda_2 \sum_{j=1}^n w_j^2$$

- Balances sparsity and small weights.
- Effect on Model Training
  - Regularization penalizes large weights, reducing model complexity.
  - The model balances fitting training data and keeping weights small.
  - Helps prevent overfitting and improves generalization.
- 🔧 How to Choose λ (Regularization Strength)?
  - Too large λ: Underfitting (model too simple)
  - Too small λ: Overfitting (model too complex)
  - Use cross-validation to find optimal  $\lambda$ .

# Example (Ridge Regression in Python): from sklearn.linear\_model import Ridge model = Ridge(alpha=1.0) # alpha = $\lambda$ model.fit(X\_train, y\_train)



- Q: Why does regularization help prevent overfitting?
  A: It constrains weights, stopping the model from fitting noise or overly complex patterns.
- Q: What's the difference between L1 and L2? A: L1 can shrink weights to zero (feature selection), L2 shrinks weights evenly but keeps them non-zero.
- Q: Can you use regularization with neural networks?

  A: Yes! Techniques like weight decay (L2 regularization), dropout, and early stopping are popular.

#### **PCA (Principal Component Analysis)**

What is PCA?

Principal Component Analysis (PCA) is an unsupervised dimensionality reduction technique used to:

- Reduce the number of features (dimensions) in data, while
- Preserving as much variance (information) as possible.

It transforms the original features into a new set of uncorrelated variables called **principal components**.

# Why Use PCA?

- High-dimensional data can be hard to visualize and process.
- Reducing dimensionality helps:
  - o Simplify models and reduce overfitting
  - o Speed up computation
  - Visualize data in 2D or 3D plots
  - o Remove redundant, correlated features

# How Does PCA Work?

- 1. Center the data by subtracting the mean of each feature.
- 2. Compute the covariance matrix of the data.
- 3. Calculate the eigenvectors and eigenvalues of the covariance matrix.
- 4. Sort eigenvectors by descending eigenvalues (variance explained).
- 5. Select the top k eigenvectors to form the principal components.
- 6. Project the original data onto these components, creating a reduced dataset.

#### **Key Terms**

Eigenvectors: directions (axes) along which the data varies most.

- Eigenvalues: magnitude of variance along each eigenvector.
- Principal Components: new feature axes (uncorrelated) that explain variance.

#### Example

- Original data: 10 features (dimensions)
- After PCA: reduce to 2 or 3 principal components while retaining ~90% of variance.
- Result: simpler dataset easier to visualize and model.

#### 

- PCA finds the directions where data spreads out the most.
- It rotates and re-scales the coordinate system to align with these directions.
- First principal component = direction of greatest variance.
- Second principal component = direction of next greatest variance, orthogonal to the first.

# PCA in Python (using sklearn):

from sklearn.decomposition import PCA

pca = PCA(n\_components=2) # reduce to 2 dimensions

X\_reduced = pca.fit\_transform(X)

# Interview Nuggets

• Q: Does PCA depend on target labels? A: No, PCA is unsupervised; it doesn't use output labels.

- Q: When should you normalize data before PCA?
- A: If features have different units or scales, normalize so PCA isn't biased toward large-scale features.
  - Q: What's the difference between PCA and feature selection?
- A: PCA creates new features (linear combinations), while feature selection picks a subset of original features.

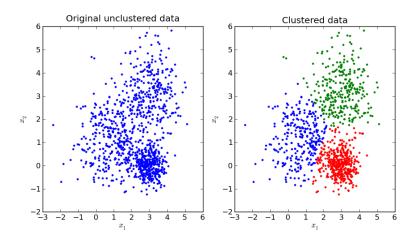
#### Summary

- PCA reduces dimensionality by projecting data onto principal components.
- Preserves maximum variance and removes feature correlation.
- Useful for visualization, speed-up, and noise reduction.
- Based on eigenvectors and eigenvalues of the covariance matrix.

#### Clustering, k-means clustering, k-nearest neighbours

# **₩** What is Clustering?

• Clustering is an **unsupervised** learning technique that groups data points into clusters so that points in the same cluster are more similar to each other than to those in other clusters.



• It's used when labels are not available and we want to find natural groupings or patterns in data.

<i>P</i> Key	Characteristics of Clustering
•	No target/output variable — it's unsupervised.
•	Similarity is usually measured by distance metrics (e.g., Euclidean distance).
•	Applications: customer segmentation, image segmentation, anomaly detection, and more.
K-Me	eans Clustering
•	One of the most popular partition-based clustering algorithms.
• points).	Groups data into K clusters, where each data point belongs to the cluster with the <b>nearest centroid</b> (mean of
How K-N	Means Works:
1.	Initialize K centroids randomly or by some heuristic.
2.	Assign each data point to the nearest centroid (cluster assignment).
3.	Recalculate centroids as the mean of assigned points.
4.	Repeat steps 2-3 until centroids stabilize or max iterations reached.

# Important Details:

- Distance metric: usually Euclidean distance.
- Sensitive to initialization can converge to local minima.

•	Requires predefining K (number of clusters).
•	Fast and efficient on large datasets.
Exampl	le Use Case:
• spende	Grouping customers based on purchasing behavior into 3 clusters: high spenders, moderate, and low rs.
🠾 K-N	earest Neighbors (KNN)
What is	KNN?
•	A <b>supervised</b> learning algorithm used for classification and regression.
•	Predicts the label/value of a data point based on the labels/values of its K closest neighbors.
How KI	NN Works:
1.	Choose K (number of neighbors).
2.	Compute the distance between the query point and all points in training data.
3.	Select the K closest points.
4.	For classification: Assign the class most common among the neighbors (majority vote).
5.	For regression: Compute the average value of the neighbors.
Importa	ant Details:
•	Distance metrics can be Euclidean, Manhattan, or others.

- No training phase it's a lazy learner.
- Sensitive to feature scaling (normalization often required).
- Works well for small to medium datasets.

# Summary

Method	Туре	Use Case	Key Points
Clustering Unsupervised		Group unlabeled data	Find natural clusters based on similarity
K-Means	Unsupervised	Partition data into K clusters	Iterative centroid update, requires K
K-Nearest Neighbors (KNN)	Supervised	Classification/Regression	Predict based on neighbors, lazy learner

Python Examples

K-Means Clustering:

from sklearn.cluster import KMeans

kmeans = KMeans(n\_clusters=3)

kmeans.fit(X)

labels = kmeans.labels\_

#### KNN Classification:

from sklearn.neighbors import KNeighborsClassifier

knn = KNeighborsClassifier(n\_neighbors=5)

knn.fit(X\_train, y\_train)

predictions = knn.predict(X\_test)

# Interview Nuggets

• Q: How to choose K in K-Means?

A: Use the elbow method (plot SSE vs. K) or domain knowledge.

Q: Difference between K-Means and KNN?

A: K-Means clusters unlabeled data (unsupervised), KNN classifies or predicts labels (supervised).

• Q: What are KNN limitations?

A: Slow on large datasets, sensitive to irrelevant features and scaling.

#### **Decision and classification trees**

- Decision Trees Overview
  - Decision Trees are a type of **supervised** learning algorithm used for both classification and regression tasks.
  - They model decisions and their possible consequences as a tree structure:
    - o Nodes represent tests on features.
    - o Branches represent outcomes of tests.
    - Leaves represent final predictions (classes or continuous values).

How Decision Trees Work

- 1. Start with the entire dataset at the root node.
- 2. Choose the best feature and threshold to split the data into subsets that are more "pure" or homogeneous.
- 3. Recursively split each subset until:
  - The subset is pure (all the same class for classification).
  - o A stopping criterion is met (max depth, minimum samples).
- 4. Assign the majority class (classification) or average value (regression) to leaf nodes.

#### **Splitting Criteria**

- For Classification Trees:
  - Gini Impurity: Measures how often a randomly chosen element would be incorrectly labeled if randomly labeled according to the distribution of labels in the subset.
- Gini=1- $\Sigma$ i=1Cpi2Gini = 1 \sum\_{i=1}^C p\_i^2Gini=1-i=1 $\Sigma$ Cpi2 where pip\_ipi is the proportion of class iii in the subset.
  - o Entropy (Information Gain):
- Entropy= $-\Sigma i=1$ Cpilog2piEntropy =  $-\sum i=1$ Cpilog2pi The split that reduces entropy the most is chosen.
  - For Regression Trees:
    - o Use metrics like Mean Squared Error (MSE) or Mean Absolute Error (MAE) to find splits.

# Classification Trees

• A classification tree is a decision tree used specifically for categorical output (class labels).

•	Final prediction a	at a leaf node is usuall	v the maiority class o	f samples in that leaf.

# \* Advantages

- Intuitive and easy to visualize.
- Requires little data preprocessing (no scaling needed).
- Can handle both numerical and categorical features.
- Captures nonlinear relationships and feature interactions.

# Limitations

- Can overfit if trees grow too deep.
- Sensitive to small changes in data.
- Greedy splitting may not find a global optimal tree.
- Pruning or ensemble methods (Random Forests, Gradient Boosted Trees) are used to improve performance.

# Example in Python (Classification Tree):

 $from \ sklearn.tree \ import \ Decision Tree Classifier$ 

clf = DecisionTreeClassifier(max\_depth=5)

clf.fit(X\_train, y\_train)

predictions = clf.predict(X\_test)

# Interview Nuggets

• Q: What is overfitting in decision trees and how to prevent it?

A: Overfitting occurs when the tree fits noise in the training data. Prevent with pruning, max depth limits, or ensemble methods.

• Q: Difference between Gini impurity and entropy?

A: Both measure impurity; Gini is slightly faster, entropy more theoretically grounded in information theory.

• Q: How do decision trees handle missing data?

A: Some implementations can handle missing values by surrogate splits or by ignoring missing features during splitting.

#### One-Hot Encoding, Label Encoding, Target Encoding, and K-Fold Cross-Validation

# 3 One-Hot Encoding

- Transforms categorical variables into a binary vector where:
  - o Each category becomes a separate feature (column).
  - The presence of a category is marked with a 1, absence with 0.
- Example: For a "Color" feature with categories Red, Green, Blue:

Color	Red	Green	Blue
Red	1	0	0
Blue	0	0	1

• Used for nominal categorical variables with no ordinal relationship.

Prevents algorithms from interpreting categories as ordinal numbers.

# Label Encoding

- Converts categories to integer labels.
- Example: Red  $\rightarrow$  0, Green  $\rightarrow$  1, Blue  $\rightarrow$  2.
- Suitable for ordinal categorical variables (where order matters).
- Warning: Some algorithms might misinterpret these integers as numeric values implying order.

# Target Encoding (Mean Encoding)

- Replaces each category with the mean of the target variable for that category.
- Useful for categorical variables in supervised learning.
- Example: If the target is 0/1 for classification, a category is replaced with the proportion of positive cases in that category.
  - Can leak target information → requires careful handling (like using K-fold target encoding).

# S K-Fold Cross-Validation

- A technique to evaluate model performance more reliably.
- Split the dataset into K equal-sized folds.
- Train on K-1 folds and test on the remaining fold.

- Repeat the process K times, each time with a different test fold.
- Final performance is averaged over all folds.
- Helps reduce variance due to random train-test splits.

# Example Python snippets:

#### One-Hot Encoding

from sklearn.preprocessing import OneHotEncoder

encoder = OneHotEncoder(sparse=False)

X\_encoded = encoder.fit\_transform(X\_categorical)

#### **Label Encoding**

from sklearn.preprocessing import LabelEncoder

encoder = LabelEncoder()

y\_encoded = encoder.fit\_transform(y\_categorical)

#### K-Fold Cross Validation

from sklearn.model\_selection import KFold

kf = KFold(n\_splits=5)

for train\_index, test\_index in kf.split(X):

X\_train, X\_test = X[train\_index], X[test\_index]

y\_train, y\_test = y[train\_index], y[test\_index]



• Q: When to use One-Hot vs Label Encoding?
A: Use One-Hot for nominal categories, Label Encoding for ordinal.

• Q: Why be careful with Target Encoding? A: It can cause target leakage if not done properly.

• Q: Why use K-Fold Cross-Validation?

A: To get more reliable, less biased performance estimates.

#### **Random Forest**

### What is Random Forest?

- Random Forest is an ensemble learning method for classification and regression.
- It builds many decision trees (a "forest") during training.
- The final prediction is made by:
  - Majority voting (classification)
  - Averaging (regression)

# Why Random Forest?

- Combats overfitting of single decision trees by averaging multiple trees.
- Improves accuracy and robustness.
- Handles large datasets with higher dimensionality.
- Can handle both categorical and numerical features.

#### How Does Random Forest Work?

- 1. Bootstrap Sampling: Create multiple training datasets by sampling with replacement from the original data (called bagging).
  - 2. For each bootstrapped dataset, build a decision tree.
  - 3. At each split in a tree, consider only a random subset of features (feature bagging) to split on.
  - 4. Trees grow fully (or to a set max depth).
  - 5. Aggregate predictions from all trees for the final output.

# Key Concepts

- Bagging: Random sampling with replacement to build diverse trees.
- Feature randomness: Randomly select subset of features at splits to reduce correlation between trees.
- Out-of-Bag (OOB) error: Estimate model error using data points not included in the bootstrap sample.

#### Advantages

- Great accuracy without much tuning.
- Robust to noise and outliers.
- Can estimate feature importance.
- Handles missing values well.

# **A** Limitations

- Less interpretable than a single decision tree.
- Can be slower and more computationally intensive.
- May overfit on noisy datasets if trees are very deep.

# Python Example

from sklearn.ensemble import RandomForestClassifier

rf = RandomForestClassifier(n\_estimators=100, max\_depth=10, random\_state=42)

rf.fit(X\_train, y\_train)

predictions = rf.predict(X\_test)

# Interview Nuggets

• Q: How does Random Forest reduce overfitting?

A: By averaging many decorrelated trees built on different bootstrap samples and feature subsets.

Q: What is feature bagging?

A: Randomly selecting a subset of features to consider at each split, which promotes tree diversity.

• Q: How to tune Random Forest?

A: Number of trees, max depth, min samples split, max features.

# **Optimizers**

An **optimizer** is an algorithm used to **adjust the weights and biases** of your machine learning model to **minimize the loss function** during training.

#### Goal:

To find the best possible parameters  $(\theta)$  that **minimize the error (loss)** on both training and validation data.

#### What does the Optimizer do?

During training, the optimizer:

1. Calculates the gradient of the loss function with respect to the model's parameters.

- 2. Updates the parameters in the direction that reduces the loss (typically via gradient descent).
- 3. Repeats this process over many iterations (epochs).

#### Why is Optimization Hard?

Because:

- Loss surfaces are often non-convex.
- The surface may contain local minima, saddle points, flat regions, or steep cliffs.
- Your optimizer needs to navigate this terrain efficiently and robustly.

**Popular Optimizers:** 

Gradient Descent and Stochastic Gradient Descent (SGD) - will be explained soon.

**SGD** with Momentum

$$egin{aligned} v_t &= \gamma v_{t-1} + \eta 
abla_{ heta} J( heta) \ heta &= heta - v_t \end{aligned}$$

• γ: momentum term (usually ~0.9)

Pros:

- Smoother convergence
- Helps escape local minima

Adam (Adaptive Moment Estimation)

Combines Momentum and RMSProp.

$$m_t = eta_1 m_{t-1} + (1-eta_1) g_t \ v_t = eta_2 v_{t-1} + (1-eta_2) g_t^2 \ \hat{m}_t = rac{m_t}{1-eta_1^t}, \quad \hat{v}_t = rac{v_t}{1-eta_2^t} \ heta = heta - \eta \cdot rac{\hat{m}_t}{\sqrt{\hat{v}_t} + \epsilon}$$

Pros:

- Very popular in DL
- Fast convergence
- Works well out of the box

Cons:

• May generalize worse than SGD on some problems

AdamW - will be explained later.

# **Adagrad**

Adapts learning rate to parameters; larger updates for infrequent parameters.

$$G_t = G_{t-1} + g_t^2 \ heta = heta - rac{\eta}{\sqrt{G_t + \epsilon}} g_t$$

Pros:

• Good for sparse data (e.g., NLP)

Cons:

• Learning rate shrinks too much over time

#### Adadelta

Improves Adagrad by limiting accumulated past gradients.

$$heta = heta - rac{ ext{RMS}[\Delta heta]_{t-1}}{ ext{RMS}[g]_t} \cdot g_t$$

- Keeps window of updates
- No need to manually set learning rate

# **Optimizer Comparison Summary**

Optimizer	Adapts LR	Momentum	Regularization	Best Use Case
Gradient Desc.	×	×	Manual	Simple convex problems, small datasets
SGD	×	<b>X</b> / <b>V</b>	Manual	Baselines, simple models
Momentum SGD	×	<b>V</b>	Manual	Deeper networks

RMSProp	V	×	Manual	RNNs, unstable gradients
Adam	V	V	No (bad decay)	Most DL models
AdamW	V	V	V	Transformers, CNNs
Adagrad	V	×	No	Sparse data (text)
Adadelta	V	×	No	Online learning
Nadam	V	✓ (Nesterov)	No	Rarely used today

#### **Gradient Descent**

#### 1. Definition & Core Concept

Gradient Descent is an iterative optimization algorithm used to find the minimum of a function. In ML, we minimize a cost (loss)  $L(\theta)$  with respect to model parameters  $\theta$  by moving  $\theta$  in the direction of steepest descent (negative gradient):

$$heta \leftarrow heta - \eta \, 
abla_{ heta} L( heta)$$

- Θ (theta): Model parameters (e.g., weights, biases)
- $L(\theta)$ : Loss function (depends on  $\theta$ )
- $\nabla \theta L(\theta)$ : Gradient of the loss with respect to  $\theta$
- η(taη): Learning rate (controls step size)

It's like a guided tour to the bottom of a valley, iteratively moving in the direction of the steepest descent until it reaches the lowest point.

#### How it works:

- Start with a random point: The algorithm begins with a random initial set of parameters.
- Calculate the gradient: It calculates the gradient (the slope or derivative) of the cost function at the current point. The gradient indicates the direction of the steepest increase in the cost function.
- Move in the opposite direction: The algorithm takes a small step in the direction opposite to the gradient (the negative gradient). This step is controlled by a learning rate, which determines the size of each step.
- Repeat: This process of calculating the gradient and moving in the opposite direction is repeated iteratively
  until the algorithm converges to a local minimum (the lowest point) of the cost function.

#### 2. Importance in AI/ML Pipelines

- Core Trainer: Backbone of training almost all neural networks and many other models (linear/logistic regression).
  - Scalability: Works efficiently on high-dimensional parameter spaces.
  - Flexibility: Variants adapt to batch sizes, noisy gradients, and sparse data.

# 3. Key Variants & Techniques

- Batch Gradient Descent: Computes gradient over the entire training set each step—stable but can be slow on large data.
- Stochastic Gradient Descent (SGD): Updates per single example—fast, introduces noise that can help escape shallow minima.
- Mini-Batch Gradient Descent: Compromise: updates on small batches (e.g., 32–256 samples) for balance of speed and stability.
  - Momentum: Adds a velocity term to smooth updates and accelerate through shallow valleys:

$$v \leftarrow \beta v + (1 - \beta) \nabla_{\theta} L$$
  
 $\theta \leftarrow \theta - \eta v$ 

- Adaptive Methods:
  - o AdaGrad: Per-parameter learning rates that shrink over time.

- RMSProp: Exponentially decaying average of squared gradients.
- o Adam: Combines momentum & adaptive scaling—most popular default.

# 4. Advantages & Limitations

Advantages	Limitations
Simple to implement & broadly applicable	Choice of learning rate critical to convergence
Scales well to large datasets & parameters	Can get stuck in local minima or saddle points
Adaptive variants require less hyper-tuning	No guarantee of global minimum on non-convex functions
Mini-batch allows GPU acceleration	Sensitive to data normalization and loss surface shape

# 5. Real-World Applications & Examples

- Neural Network Training: All deep-learning frameworks use GD variants under the hood.
- Recommendation Systems: Matrix factorization via SGD on sparse ratings.
- Logistic Regression: Binary classifiers trained with (mini-batch) GD.
- Computer Vision: Fine-tuning pre-trained CNNs with Adam or SGD+Momentum.

# 6. Sample Code Snippet (TensorFlow/Keras Custom Training Loop)

```
import tensorflow as tf
# Simple linear model: y = Wx + b
W = tf.Variable(tf.random.normal([1]), name='weight')
b = tf.Variable(tf.zeros([1]), name='bias')
optimizer = tf.keras.optimizers.SGD(learning_rate=0.01, momentum=0.9)
@tf.function
def train_step(x, y_true):
  with tf.GradientTape() as tape:
    y_pred = W * x + b
    loss = tf.reduce_mean((y_true - y_pred)**2)
  grads = tape.gradient(loss, [W, b])
  optimizer.apply_gradients(zip(grads, [W, b]))
  return loss
# Example training loop
for epoch in range(100):
  loss = train_step(x_data, y_data)
  if epoch % 10 == 0:
    print(f"Epoch {epoch}: Loss = {loss.numpy():.4f}")
```

#### 7. Common Interview Questions & How to Answer

• Q: "How do you choose the learning rate?"

A: "Start small (e.g., 1e-3), monitor loss curve-too large causes divergence, too small makes training slow; use

learning-rate schedules or adaptive optimizers."

• Q: "Why use mini-batch over full-batch or SGD?"

A: "Mini-batches balance gradient estimate stability and computational efficiency (GPU vectorization), plus introduce helpful noise."

• Q: "What's the role of momentum?"

A: "Momentum smooths oscillations by accumulating past gradients, helping traverse ravines and accelerate convergence."

• Q: "When might Adam underperform SGD?"

A: "On large-scale vision tasks, SGD with momentum sometimes generalizes better, as Adam's adaptive steps can overfit to noise."

#### Stochastic gradient descent

→ What is Stochastic Gradient Descent (SGD)?

- SGD is an optimization algorithm used to **minimize the loss function in machine learning models**, especially neural networks.
- It's a variant of Gradient Descent that updates model parameters using only one training example (or a small batch) at a time, rather than the entire dataset.

$$heta = heta - \eta \cdot 
abla_{ heta} J( heta)$$

- θ: parameters (weights)
- η: learning rate
- $\nabla \theta J(\theta)$ : gradient of the loss
- How SGD Works:
  - Initialization: Start with initial values for the model parameters.
  - Iteration:
    - Choose a data point: Randomly select a single data point from the training set.
    - Calculate the gradient: Calculate the gradient of the loss function with respect to the model parameters for that single data point.
    - Update parameters: Update the model parameters by subtracting a scaled gradient from the current parameters. The scaling factor is a learning rate (a hyperparameter that controls the step size).

 Repeat: Repeat steps 2 and 3 for a specified number of iterations or until a convergence criterion is met.

# \* Key Points

- Faster updates and more frequent parameter tuning than batch gradient descent.
- Noisy updates because gradients are based on single samples can help escape local minima.
- Requires careful learning rate tuning to ensure convergence.
- Often combined with mini-batches (Mini-Batch Gradient Descent) for a balance of stability and speed.

# Why Use SGD?

- Works well on large datasets where full batch gradient descent is computationally expensive.
- Introduces stochasticity (randomness) that can improve generalization.
- Widely used with deep learning frameworks.

# Python snippet (using PyTorch):

```
import torch
```

import torch.optim as optim

```
model = MyModel()
```

optimizer = optim.SGD(model.parameters(), lr=0.01)

for epoch in range(num\_epochs):

for x\_batch, y\_batch in data\_loader: # data\_loader yields mini-batches

optimizer.zero\_grad()

outputs = model(x\_batch)

loss = loss\_fn(outputs, y\_batch)
loss.backward()
optimizer.step()

## Interview Nuggets

Q: Difference between batch gradient descent and SGD?

A: Batch GD uses the entire dataset for each update; SGD uses single samples, making it faster but noisier.

Q: What is a learning rate and why is it important?

A: It controls step size during updates; too high causes divergence, too low slows training.

Q: What are common SGD variants?

A: Mini-batch SGD, SGD with momentum, Adam optimizer, RMSProp.

## Vectorization

## What is Vectorization?

- Vectorization is the process of converting operations that use explicit loops into batch operations using vectors or matrices.
- Instead of processing data element-by-element in a loop, vectorized operations use efficient linear algebra routines to perform computations on entire arrays at once.
  - Commonly done using libraries like NumPy, TensorFlow, or PyTorch.

## Q Why Vectorize?

- Speed: Vectorized code runs much faster because it leverages optimized, low-level implementations (often in C/C++).
  - Simplicity: Code becomes more concise and easier to read.
  - Hardware optimization: Takes advantage of CPU SIMD instructions and GPU parallelism.

# ★ Example Non-vectoriz

Non-vectorized (loop-based) Python code:

# Element-wise addition of two lists

```
a = [1, 2, 3]
```

$$b = [4, 5, 6]$$

for i in range(len(a)):

```
c.append(a[i] + b[i])
```

Vectorized using NumPy:

import numpy as np

```
a = np.array([1, 2, 3])
```

b = np.array([4, 5, 6])

c = a + b # element-wise addition without explicit loop

# In Neural Networks

- Vectorization is essential for:
  - o Efficient matrix multiplications during forward and backward passes.
  - o Batch processing multiple samples simultaneously.
  - Leveraging GPUs for parallel computation.

# Interview Nuggets

• Q: Why is vectorization important in ML?

A: It drastically speeds up computations and enables handling large datasets efficiently.

• Q: Can you give an example of vectorized operations?

A: Matrix multiplication, element-wise addition, dot products.

• Q: What's the downside of non-vectorized code?

A: It's slower and harder to maintain, especially with large datasets.

## AdaBoost, Gradient Boosting, and XGBoost

- AdaBoost (Adaptive Boosting)
- AdaBoost is an ensemble boosting method that combines many weak learners (usually shallow decision trees called stumps) into a strong classifier.
  - Key idea:
    - o Train each weak learner sequentially.
    - After each learner, increase the weight of misclassified samples so the next learner focuses more on the hard cases.
  - Final prediction is a weighted vote of all learners.
  - Good for reducing bias, but sensitive to noisy data and outliers.

## Y Gradient Boosting

- Also an ensemble boosting method that builds trees sequentially.
- Instead of reweighting samples, each new tree is trained to predict the residual errors (gradients) of the combined previous trees.

- The model is optimized via gradient descent on a loss function.
- Allows use of arbitrary differentiable loss functions (regression, classification).
- More flexible and powerful than AdaBoost.

# 

• An optimized implementation of gradient boosting.

## Features:

- o Regularization (L1 & L2) to reduce overfitting.
- o Parallel and distributed computing.
- o Handling missing data automatically.
- Tree pruning and early stopping.
- Supports custom loss functions.
- Extremely popular in ML competitions (Kaggle) for its speed and accuracy.

# Key Differences

Aspect	AdaBoost	Gradient Boosting	XGBoost
Base Learners	Weak learners (stumps)	Decision trees	Decision trees (optimized)
Training Focus	Reweight misclassified samples	Fit residual errors	Same as Gradient Boosting + optimizations
Loss Function	Exponential loss	Any differentiable loss	Any differentiable loss
Regularization	No explicit regularization	Limited	L1, L2 regularization
Speed	Slower	Moderate	Fast, highly optimized

		Handling Missing	No	No	Yes
--	--	------------------	----	----	-----

Python Example (XGBoost):

import xgboost as xgb

model = xgb.XGBClassifier(n\_estimators=100, max\_depth=5, learning\_rate=0.1)

model.fit(X\_train, y\_train)

preds = model.predict(X\_test)

# Interview Nuggets

• Q: What is boosting?

A: Sequentially combining weak learners to create a strong learner by focusing on errors.

- Q: How does Gradient Boosting differ from AdaBoost?
- A: Gradient Boosting fits residuals using gradient descent, AdaBoost reweights samples based on errors.
  - Q: Why use XGBoost over vanilla Gradient Boosting?

A: XGBoost is faster, supports regularization, missing data handling, and offers better control over training.

## **Cosine Similarity**

- What is Cosine Similarity?
  - A measure of similarity between two non-zero vectors in an inner product space.
  - It calculates the cosine of the angle between the vectors.
  - Values range from -1 to 1:

<ul> <li>1 means vectors point in the same direction</li> </ul>	on (maximum similarity).
o 0 means vectors are orthogonal (no simila	urity).
o -1 means vectors point in opposite direction	ons (maximum dissimilarity).
Formula	
For two vectors A\mathbf{A}A and B\mathbf{B}B:	
$\label{lem:cosineSimilarity} $$ CosineSimilarity(A,B)=A\cdot B \ \ A \ \ \ B \ \ = \Sigma i=1nAiBi\Sigma i=1nAi2\Sigma i \ \ \ \ \ \ \ \ \ \ \ \ \ \ \ \ \ \ $	$\{B\}\setminus = \frac{i=1}^n A_i B_i}{\sqrt{\sum_{i=1}^n}}$
Where:	
• A · B\mathbf{A} \cdot \mathbf{B}A · B = dot product	
•	// = magnitudes (Euclidean norms)
Why use Cosine Similarity?	
Measures orientation similarity regardless of magn	itude.
Commonly used in:	
o Text analysis / NLP (e.g., comparing docu	ment vectors or word embeddings).
Recommendation systems.	
o Clustering and information retrieval.	
Python Example	
python	

def cosine\_similarity(A, B):

```
dot_product = np.dot(A, B)
norm_A = np.linalg.norm(A)
norm_B = np.linalg.norm(B)
return dot_product / (norm_A * norm_B)

vec1 = np.array([1, 2, 3])
vec2 = np.array([4, 5, 6])
print(cosine_similarity(vec1, vec2)) # Output ~0.9746
```

# Interview Nuggets

- Q: Why is cosine similarity preferred over Euclidean distance in text?

  A: Because it measures angle (orientation), so document length differences don't affect similarity.
- Q: What does a cosine similarity of 0 mean? A: Vectors are orthogonal, meaning no similarity.
- Q: Can cosine similarity be negative? A: Yes, if vectors point in opposite directions.

## **Support Vector Machine**

What is SVM?

- Support Vector Machine (SVM) is a **supervised** learning algorithm mainly used for classification (and regression).
- It finds the optimal hyperplane that best separates classes by maximizing the margin the distance between the hyperplane and the nearest data points from each class.
  - These nearest points are called support vectors.

## What is a Kernel?

A kernel is a function that computes a dot product (similarity) in a high-dimensional feature space — without explicitly transforming the data into that space. This is known as the kernel trick.

## Marks:

1. Data Representation:

Data points are represented as vectors in a multi-dimensional space, where each dimension corresponds to a feature.

2. Finding the Optimal Hyperplane:

SVM aims to find the hyperplane that maximizes the margin while minimizing the classification errors.

3. Support Vectors:

The data points closest to the hyperplane are identified as support vectors.

4. Decision Boundary:

The hyperplane acts as the decision boundary, separating different classes.

## Key Concepts

- Margin: Distance between hyperplane and nearest points.
- Support Vectors: Data points that influence the position of the hyperplane.
- Kernel Trick: Transforms data to higher dimensions to handle nonlinear problems without explicit mapping.
- C parameter: Controls trade-off between margin size and classification error (soft margin).
- Simple example in Python (using scikit-learn):

from sklearn.svm import SVC

model = SVC(kernel='rbf', C=1.0, gamma='scale')

model.fit(X\_train, y\_train)

predictions = model.predict(X\_test)

# Interview Nuggets

Q: What are support vectors?

A: Critical training samples that define the decision boundary.

• Q: Why use kernels?

A: To handle data that's not linearly separable by implicitly mapping it to higher dimensions.

• Q: What does the parameter C do?

A: Balances margin maximization and misclassification tolerance.

Q: Difference between hard-margin and soft-margin SVM?

A: Hard-margin requires perfect separation; soft-margin allows some misclassification for better generalization.

## The Essential Main Ideas of Neural Networks

## 1. Definition & Core Concept

A Neural Network is a computational model inspired by the human brain's network of neurons. It consists of layers of interconnected nodes (neurons) that transform input data through weighted connections and nonlinear activation functions to learn representations and solve tasks.

## 2. Importance in AI/ML Pipelines

- $\bullet$  Feature Learning: Automatically extracts hierarchical features from raw data (e.g., pixels  $\rightarrow$  edges  $\rightarrow$  objects).
  - Universal Function Approximation: Can approximate any continuous function given sufficient capacity.
- End-to-End Learning: Permits seamless training from raw input all the way to output, reducing manual feature engineering.

## 3. Key Components:

## Neurons (Nodes):

The basic processing units in a neural network. They receive input, apply a mathematical function (activation function), and produce an output.

## Layers:

Neurons are organized into layers. Input layers receive the initial data, hidden layers process the information, and output layers produce the final result.

#### **Connections and Weights:**

Neurons are connected, and each connection has an associated weight, representing the strength of the connection. These weights are adjusted during training.

## **Activation Functions:**

Non-linear functions that determine the output of a neuron based on its input.

## 4. How Neural Networks Learn:

- Training Data: Neural networks are trained on large datasets of labeled examples.
- **Learning Algorithm:** An algorithm (like gradient descent) adjusts the weights based on the difference between the network's predictions and the actual labels.
- Iteration: The process of training involves repeatedly feeding data to the network, making predictions, comparing them to the actual labels, and adjusting the weights until the network achieves a desired level of accuracy.

## 5. Key Architectures & Techniques

- Feedforward (MLP): Simple layered structure; data moves one direction.
- Convolutional Nets (CNNs): Leverages spatial locality with convolutional filters for images.
- Recurrent Nets (RNNs/LSTMs/GRUs): Captures sequential dependencies—text, time series.
- Transformers: Attention-based, excels in language tasks (e.g., GPT, BERT).
- Training Tricks:
  - o Backpropagation for gradient computation
  - o Batch Normalization, Dropout for regularization
  - Adaptive Optimizers: Adam, RMSProp

Learning Rate Schedules (warmup, decay)

# 6. Advantages & Limitations

Advantages	Limitations
Learns complex, non-linear relationships	Requires large labeled datasets
Highly adaptable across domains	Prone to overfitting without regularization
Scales well with compute & data	Often "black box" – hard to interpret
State-of-the-art performance in vision/NLP	High computational cost & energy usage

# 7. Real-World Applications & Examples

- Computer Vision: Image classification (ResNet, EfficientNet), object detection (YOLO, Faster R-CNN).
- Natural Language Processing: Language modeling and translation (Transformer, GPT series).
- Healthcare: Disease diagnosis from medical scans.
- Autonomous Vehicles: Sensor fusion and scene understanding.

• Finance: Fraud detection, algorithmic trading.

```
8. Sample Code Snippet (TensorFlow/Keras)
```

```
from tensorflow.keras import layers, models
def build_simple_nn(input_shape, num_classes):
  model = models.Sequential([
    layers.Input(shape=input_shape),
    layers.Dense(128, activation='relu'),
    layers.Dropout(0.3),
    layers.Dense(64, activation='relu'),
    layers.Dense(num_classes, activation='softmax')
  ])
  model.compile(optimizer='adam',
          loss='sparse_categorical_crossentropy',
          metrics=['accuracy'])
  return model
# Example usage:
# model = build_simple_nn((784,), 10)
# model.fit(x_train, y_train, epochs=10, batch_size=32)
```

- 9. Common Interview Questions & How to Answer
  - Q: "Explain how backpropagation works."

A: "Backpropagation computes gradients of the loss w.r.t. each weight by applying the chain rule layer by layer, then updates weights via gradient descent."

• Q: "Why do we need activation functions?"

A: "They introduce non-linearity, allowing networks to learn complex mappings beyond linear transformations."

Q: "How would you prevent overfitting?"

A: "Use techniques like dropout, L1/L2 regularization, data augmentation, and early stopping."

• Q: "When to choose a CNN vs. a Transformer?"

A: "CNNs excel on grid-structured data (images); Transformers on sequence data with long-range dependencies (text)."

## Train, Validate, Test

## 1. Training Set

- The portion of the dataset used to fit the model i.e., to learn the model parameters (weights, biases).
- The model "sees" this data during training.

## 2. Validation Set

- A separate subset used during training to tune hyperparameters and evaluate model performance on unseen data.
  - Helps detect overfitting and guides choices like early stopping, model architecture, learning rate.
  - Model does not learn from this data, only used for evaluation.

## 3. Test Set

- A final, independent dataset used after training to assess how well the model generalizes to new, unseen data.
  - Provides an unbiased evaluation of model performance.

## Why Split?

• To avoid overfitting: without separate validation/test sets, the model might just memorize training data.

To ensure the model works well on real-world data (generalization).



# Typical Split Ratios

Dataset	Common Percentage
Training	70-80%
Validation	10-15%
Test	10-15%

# Interview Nuggets

• Q: Why not test on training data?

A: It biases evaluation; model accuracy will be unrealistically high.

- Q: What's the difference between validation and test sets? A: Validation is for tuning during training; test is for final unbiased evaluation.
  - Q: What if dataset is small?

A: Use techniques like cross-validation to maximize data use.

## **Activation Functions**

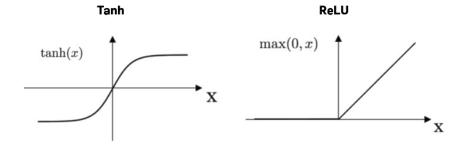
## What are Activation Functions?

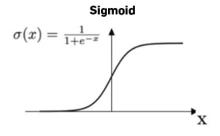
- Functions applied to a neuron's output in a neural network.
- Introduce non-linearity so the network can learn complex patterns.

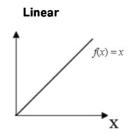
• Without activation functions, neural networks would just be linear models, no matter how many layers.

# Common Activation Functions

Activation	Formula	Properties & Use Cases
Sigmoid	$\sigma(x)=11+e-x \cdot \sin(x) = \frac{1}{1+e^{-x}}\sigma(x)=1+e-x1$	Output between 0 and 1; used in binary classification output layers. Saturates at extremes causing vanishing gradients.
Tanh	tanh(x)=ex-e-xex+e-x\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}tanh(x)=ex+e-xex-e-x	Output between -1 and 1; zero-centered, better than sigmoid but still suffers from vanishing gradients.
ReLU	$ReLU(x)=max(0,x)\setminus \{ReLU\}(x) = \max(0,x)$ $x)ReLU(x)=max(0,x)$	Most popular; simple and efficient; mitigates vanishing gradient. Can "die" if neurons output zero consistently.
Leaky ReLU	LeakyReLU(x)=max(0.01x,x)\text{LeakyReLU}(x) = \max(0.01x, x)LeakyReLU(x)=max(0.01x,x)	Variant of ReLU allowing small gradient when x<0x < 0x<0, reducing "dead neuron" problem.
Softmax	softmax(xi)=exi∑jexj\text{softmax}(x_i) = \frac{e^{x_i}}{\sum_j e^{x_j}}softmax(xi)=∑jexjexi	Converts logits to probabilities; used in multi-class classification output layers.







Activation functions graphs

# Why Activation Functions Matter

- Enable the network to learn non-linear decision boundaries.
- Affect how gradients propagate during training.
- Choosing the right activation can impact training speed and model accuracy.

# Interview Nuggets

• Q: Why not use only linear activation?

A: Multiple linear layers collapse into one linear transformation — no complex patterns learned.

- Q: What problems does ReLU solve compared to sigmoid?
- A: Reduces vanishing gradient problem and speeds up training.
  - Q: When to use Softmax?

A: At the output layer for multi-class classification to get probability distribution.

## **Loss functions**

What is a Loss Function?

A loss function measures the difference between the model's predictions and the actual target values. It's a key ingredient in training: it tells the model how wrong it is and guides how to update the weights.

# Nole in Learning

During training, the model tries to minimize the loss using an optimization algorithm like gradient descent.

Smaller loss = better performance (on training data at least).

Task	Loss Function	Formula / Description	Use Case
Regression Mean Squared Error (MSE)		1n∑(yi-y^i)2\frac{1}{n} \sum (y_i - \hat{y}_i)^2n1∑(yi-y^i)2	Penalizes large errors more
	Mean Absolute Error (MAE)	( \frac{1}{n} \sum	y_i - \hat{y}_i
	Huber Loss	Combines MSE & MAE: quadratic for small errors, linear for large	Robust regression
	Log-Cosh Loss	$ \begin{split} & \Sigma log(cosh(y^i-yi)) \setminus sum \cdot log( \cdot (hat{y}_i - y_i)) & \Sigma log(cosh(y^i-yi)) \end{split} $	Smooth alternative to MSE
	Quantile Loss	max(q(y-y^),(q-1)(y-y^))\max(q(y - \hat{y}), (q - 1)(y - \hat{y}))max(q(y-y^),(q-1)(y-y^))	Predicting quantiles
	Poisson Loss	y^-ylog(y^)\hat{y} - y \log(\hat{y})y^-ylog(y^)	Count data regression

Tweedie Loss	Models compound Poisson-Gamma distributions	Insurance & actuarial modeling
MAPE (Mean Absolute Percentage Error)	(\frac{100%}{n}\sum\left	y_i - \hat{y}_i}{y_i} \right

Binary Classification	Loss Function	Formula / Description	Use Case
	Binary Cross-Entropy	-[ylog(y^)+(1-y)log(1-y^)]-[y \log(\hat{y}) + (1 - y) \log(1 - \hat{y})]-[ylog(y^)+(1-y)log(1-y^)]	Binary classification
	Hinge Loss	max(0,1-y·y^)\max(0, 1 - y \cdot \hat{y})max(0,1-y·y^)	SVM binary classification
	Focal Loss	-α(1-y^)γylog(y^)-\alpha(1 - \hat{y})^\gamma y \log(\hat{y})-α(1-y^)γylog(y^)	Class imbalance, e.g. detection
	Log Loss	Alias of binary cross-entropy	Logistic regression

Multi-Class Classification	Loss Function	Formula / Description	Use Case
	Categorical Cross-Entropy	-∑yilog(y^i)-\sum y_i \log(\hat{y}_i)-∑yilog(y^i)	Multi-class with one-hot labels
	Sparse Categorical Cross-Entropy	Like above, but uses integer labels	Efficient multi-class
	Kullback-Leibler Divergence (KLDiv)	Σyilog(yiy^i)\sum y_i \log\left(\frac{y_i}{\hat{y}_i} \right)Σyilog(y^iyi)	Comparing distributions

Label Smoothing Loss	Modified cross-entropy with smoothed labels	Regularization, prevent overconfident models
Focal Loss (Multi-class)	Variant of cross-entropy focusing on hard examples	Object detection (e.g., RetinaNet)

Ranking / Margin-Based	Loss Function	Formula / Description	Use Case
	Hinge Loss	max(0,1-y·y^)\max(0, 1 - y \cdot \hat{y})max(0,1-y·y^)	SVMs
	Squared Hinge Loss	(max(0,1-y·y^))2(\max(0, 1 - y \cdot \hat{y}))^2(max(0,1-y·y^))2	SVM variant
	Triplet Loss	( \max(0,	
	Contrastive Loss	Penalizes similar pairs with large distance and dissimilar pairs with small distance	Siamese networks
	Ranking Loss / Pairwise		Learning to rank

Generative Models / Probabilistic	Loss Function	Formula / Description	Use Case
	Negative Log-Likelihood	( -\log P(y	\hat{y}) )
	ELBO (Evidence Lower Bound)	Used in Variational Autoencoders (VAEs) combining KL divergence + NLL	Generative modeling
	Adversarial Loss	From GANs: generator loss vs. discriminator loss	GAN training

# **Evaluation Metrics**

Core Evaluation Metrics by Task

# Classification Metrics

Metric	Formula / Meaning	Use Case
Accuracy	TP + TNTP + FP + TN + FN\frac{\text{TP + TN}}{\text{TP + FP + TN + FN}}TP + FP + TN + FNTP + TN	Good for balanced classes
Precision	TPTP + FP\frac{\text{TP}}{\text{TP + FP}}TP + FPTP  → "How many predicted positives were correct?"	Important when false positives are costly (e.g., spam detection)
Recall (Sensitivity)	TPTP + FN\frac{\text{TP}}{\text{TP + FN}}TP + FNTP  → "How many actual positives were found?"	Important when false negatives are costly (e.g., disease diagnosis)

F1 Score	2 · Precision · RecallPrecision + Recall2 \cdot \frac{\text{Precision} \cdot \text{Recall}}{\text{Precision + Recall}}2 · Precision + RecallPrecision · Recall	Balances precision and recall, especially for imbalance
Specificity	TNTN + FP\frac{\text{TN}}{\text{TN + FP}}TN + FPTN	True Negative Rate
ROC AUC	Area under ROC curve (TPR vs. FPR)	Overall model ranking ability across thresholds
PR AUC	Area under Precision-Recall curve	Especially good for imbalanced datasets
Log Loss	Cross-entropy loss; penalizes wrong confident predictions	Classifier probability output quality
Cohen's Kappa	Agreement vs. chance agreement	Classification reliability metric
Matthews Corr. Coefficient (MCC)	Balanced metric even for imbalanced classes	Correlation between true and predicted classes

✓ TP = True Positive, TN = True Negative, FP = False Positive, FN = False Negative

# Regression Metrics

Metric	Formula / Description	Use Case
MSE (Mean Squared Error)	1n∑(yi-y^i)2\frac{1}{n} \sum (y_i - \hat{y}_i)^2	Penalizes large errors heavily
RMSE (Root MSE)	MSE\sqrt{MSE}	Same as MSE but interpretable in output units
MAE (Mean Absolute Error)	( \frac{1}{n} \sum	y_i - \hat{y}_i

R² (Coefficient of Determination)	Proportion of variance explained by the model	Ranges from −∞-\infty to 1
Adjusted R <sup>2</sup>	Penalizes adding unnecessary variables	Multiple regression
MAPE (Mean Absolute Percentage Error)	% difference averaged over all points	Time series & business forecasting
Huber Loss	Combines MSE and MAE: less sensitive to outliers	Robust regression

# Clustering Metrics

Metric	Description	Use Case
Silhouette Score	How close samples are to their own cluster vs. other clusters	Internal cluster quality
Davies-Bouldin Index	Lower is better; compares intra- and inter-cluster distance	Cluster separation
Adjusted Rand Index (ARI)	Measures similarity between true labels and cluster assignments	With ground truth
Normalized Mutual Info (NMI)	Mutual information scaled to [0,1]	Label comparison

# Ranking & Recommendation Metrics

Metric	Description	Use Case
MAP (Mean Average Precision)	Average of precision at each relevant item	Ranking problems

NDCG (Normalized Discounted Cumulative Gain)	Emphasizes order of relevant items	Recommender systems
Hit Rate / Recall@K	Whether relevant item is in top-K predictions	Recommendations

# Time Series Metrics

Metric	Description
MASE	Mean Absolute Scaled Error — normalized version of MAE
SMAPE	Symmetric MAPE
RMSE / MAE / MAPE	Used as in regression, applied to time series

# ↑ Choosing the Right Metric

Scenario	Preferred Metric(s)
Imbalanced classification	F1, ROC AUC, PR AUC, MCC
Multi-class classification	Accuracy + macro/micro-averaged F1
Regression with outliers	MAE, Huber Loss
High cost of false negatives	Recall, F1
Recommendation systems	NDCG, MAP, Hit Rate@K

Time series forecasting	MAPE, RMSE, SMAPE
-------------------------	-------------------

## The Chain Rule

## 1. Definition & Core Concept

The Chain Rule is a fundamental theorem in calculus for computing the derivative of a composition of functions.

If y=f(g(x))y = f(g(x))y=f(g(x)), then

 $dydx=f'(g(x))\times g'(x).\frac\{dy\}\{dx\}=f'\wedge (g(x))\wedge (g(x))\wedge (g(x))\otimes (g(x)\otimes (g(x))\otimes (g(x))\otimes (g(x)\otimes (g(x))\otimes (g(x)\otimes (g(x))\otimes (g(x)\otimes (g($ 

## 2. Importance in AI/ML Pipelines

- **Backpropagation**: Underpins how neural networks learn—gradients flow layer-by-layer via repeated application of the chain rule.
- Auto-diff Engines: Frameworks like TensorFlow and PyTorch rely on chain-rule logic to build computation graphs and compute derivatives automatically.
- **Optimization**: Any gradient-based optimizer (SGD, Adam) depends on correct chain rule application to update model parameters.

## 3. Key Techniques & Variants

- $\bullet$  Forward Mode Differentiation: Propagates derivatives from inputs to outputs—efficient when #inputs  $\ll$  #outputs.
- Reverse Mode Differentiation (**Backprop**): Propagates gradients from output back to inputs—efficient when #outputs  $\ll$  #inputs (common in ML).
  - Symbolic Differentiation: Algebraically applies chain rule to symbolic expressions (e.g., in SymPy).
- Numerical Differentiation: Approximates derivatives via finite differences—useful for sanity-checks but less precise/efficient.

## 4. Advantages & Limitations

Advantages	Limitations
Exact gradient computation (auto-diff)	Can incur memory overhead storing intermediate activations (reverse mode)
Scales to arbitrarily deep compositions	Numerical instability (vanishing/exploding gradients)
Enables efficient optimization	Requires careful graph-structure handling

# 5. Real-World Applications & Examples

- Training Deep Networks: Every layer's weight update uses chain rule to compute ∂Loss/∂Weight.
- Physics Simulations: Sensitivity analysis of complex models.
- Financial Modeling: Greeks (sensitivities) in option pricing rely on chain rule for composite payoff functions.

# 6. Sample Code Snippet (PyTorch Autograd)

import torch

# Define scalar input with gradient tracking

x = torch.tensor(2.0, requires\_grad=True)

$$#y = f(g(x)) = (3x^2 + 1)^2$$

$$y = (3 * x**2 + 1)**2$$

# Compute derivative dy/dx via backprop

y.backward()

```
print(f"x = {x.item()}, y = {y.item()}, dy/dx = {x.grad.item()}")

# Expected: y = (3*4 + 1)^2 = 13^2 = 169; dy/dx = 2*(3x^2 + 1)*6x = 2*13*6*2 = 312
```

## 7. Common Interview Questions & How to Answer

- Q: "Derive the gradient of y=sin(5x3)y =  $\sin(5x^3)y=\sin(5x^3)$ ." A: dydx=cos(5x3)×15x2\frac{dy}{dx} =  $\cos(5x^3)$ \times 15x^2dxdy=cos(5x3)×15x2.
- Q: "Why do deep networks suffer vanishing gradients?"
  A: "Repeated chain-rule multiplications of derivatives <1 (e.g., sigmoid'<1) cause gradients to shrink exponentially."
  - Q: "Explain forward vs. reverse mode auto-diff."

A: "Forward mode propagates derivative of each sub-expression alongside its value (good for few inputs); reverse mode accumulates sensitivities backwards from outputs (good for scalar loss functions)."

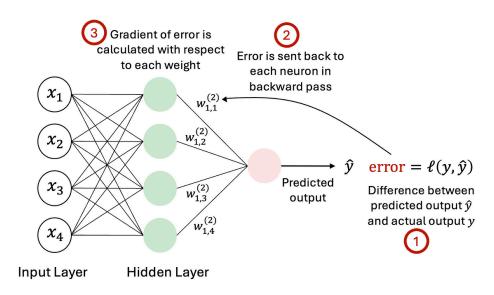
• Q: "How does the chain rule enable backpropagation?"

A: "Backprop applies the chain rule at each layer: gradient of loss w.r.t. weights = gradient of loss w.r.t. layer output × derivative of output w.r.t. Weights."

## **Backpropagation Main Ideas**

## 1. Definition & Core Concept

Backpropagation is the algorithm for efficiently computing the gradient of a neural network's loss with respect to each weight. It applies the chain rule layer by layer in reverse—from the output back to the inputs—allowing weight updates based on how much each contributed to the error.



# 2. Why It Matters in AI/ML

- Foundation of Learning: Without backprop, training deep networks would be infeasible.
- Scalability: Handles millions of parameters by reusing intermediate computations.

• Auto-Diff Engines: Modern frameworks (TensorFlow, PyTorch) implement backprop under the hood.

## 3. Step-by-Step Process

## 1. Forward Pass

 $\circ$  Compute activations a(I)a^{(I)}a(I) and pre-activations z(I)z^{(I)}z(I) for each layer III.

## 2. Compute Loss

o Compare network output to ground truth via loss function LLL.

## 3. Backward Pass

- $\begin{array}{ll} \circ & \text{Output Layer:} \\ & \delta(L) = \nabla aL \circ \sigma'(z(L)) \setminus (L) = \lambda_{a}L \cdot (z^{(L)} \cdot (L)) \\ \end{array}$
- $\begin{array}{ll} \circ & \text{Hidden Layers (for I=L-1I=L-1 down to 1):} \\ \delta(I)=(W(I+1)T\ \delta(I+1)) \circ \sigma'(z(I)) \ \ \text{delta^{(I)}} = \\ \text{bigl(W^{(I+1)T},\delta^{(I+1)}\circ\sigma'(z(I))} \end{array}$

## 4. Gradient Computation

- $\qquad \text{Weight gradients: } \nabla W(I)L = \delta(I) \ a(I-1)T \cap abla_{W^{(I)}}L = \delta(I)A^{(I)} \cap A^{(I-1)T} \nabla W(I)L = \delta(I)A(I-1)T \cap A^{(I)} \cap A^{(I$
- Bias gradients:  $\nabla b(I)L=\delta(I) \cap \{b^{(I)}\}L = \det^{(I)} \nabla b(I)L=\delta(I)$

# 5. Parameter Update

e.g., via Gradient Descent:  $W(I) \leftarrow W(I) - \eta \nabla W(I) L, b(I) \leftarrow b(I) - \eta \nabla b(I) L W^{(I)} \setminus W^{(I)} - \beta W^{(I)} L, \gamma b(I) + \beta^{(I)} L W^{(I)} + \beta^$ 

## 4. Advantages & Caveats

Advantages	Caveats
Reuses computed activations for efficiency	Can suffer vanishing/exploding gradients

Enables deep models to be trained end-to-end	Sensitive to choice of activation & init
Works with any differentiable activation	Requires computing and storing all z(l),a(l)z^{(l)},a^{(l)}z(l),a(l)

# 5. Real-World Example & Code (PyTorch)

```
import torch
import torch.nn as nn
import torch.optim as optim
# Simple 2-layer network
model = nn.Sequential(
  nn.Linear(10, 20),
  nn.ReLU(),
  nn.Linear(20, 1)
criterion = nn.MSELoss()
optimizer = optim.SGD(model.parameters(), lr=0.01)
# Dummy data
x = torch.randn(5, 10)
y = torch.randn(5, 1)
# Forward
pred = model(x)
```

```
loss = criterion(pred, y)

# Backprop + update

optimizer.zero_grad()

loss.backward() # ← automatic backprop

optimizer.step()
```

- 6. Common Interview Questions & Answers
  - Q: "Why do we apply the chain rule in backprop?"

A: "To decompose the overall gradient into products of local derivatives at each layer, making computation tractable."

• Q: "How do activations affect gradient flow?"

A: "Functions like sigmoid can shrink gradients (< 1) causing vanishing; ReLU avoids this by having derivative 1 for positive inputs."

• Q: "What's the difference between automatic and manual backprop?"

A: "Auto-diff tracks operations on tensors to build a computation graph and then applies backprop; manual means you derive and implement each gradient yourself."

• Q: "How can you mitigate vanishing gradients?"

A: "Use ReLU/LeakyReLU, proper weight initialization (e.g. He), batch normalization, or skip connections (ResNets)."

## **Hyperparameters**

Hyperparameters are settings or configurations that control the learning process of a model and are set before training begins.

## **Learning Rate**

- Controls how much the model's weights are updated during training.
- A small LR → slow learning, may get stuck in local minima.
- A large LR → faster learning but risk overshooting minima or divergence.
- Often requires tuning and sometimes decayed during training.

## **Number of Layers**

- Refers to how many layers your neural network has.
- More layers = deeper network → can learn more complex patterns.
- Too many layers → risk of overfitting and training difficulties (vanishing gradients).
- Common deep nets: 3-100+ layers depending on task and architecture.

## **Batch Size**

- Number of training samples used to compute one update of model weights.
- Small batch size (e.g., 16, 32):
  - $\circ \qquad \text{Noisier gradient estimate} \rightarrow \text{can improve generalization}.$
  - Slower training per epoch (more updates).
- Large batch size (e.g., 128, 256, 512):
  - o More stable gradient.
  - Faster training but can hurt generalization.
- Also limited by hardware memory.

## **Number of Epochs**

- How many complete passes over the training dataset.
- Too few → underfitting.
- Too many → overfitting.
- Use early stopping to avoid overtraining.

## **Optimizer**

- Algorithm that updates model weights based on gradients.
- Examples:
  - o SGD (Stochastic Gradient Descent) simple and effective.
  - Adam adaptive learning rates, widely used.
  - o RMSprop, Adagrad etc.

## **Activation Functions**

<ul> <li>Non-linear functions applied to la</li> </ul>	yers.
--	-------

•	Common	choices:	Rel U	Siamoid.	Tanh.	Leaky	Rel U

# **Dropout Rate**

- Fraction of neurons randomly ignored during training to prevent overfitting.
- Typical values: 0.1-0.5.

## **Weight Initialization**

- How weights are set before training.
- Good initialization speeds convergence and prevents problems.

## **Kernel Size**

• Used in CNN training, refers to the dimensions of the filter that slides across the input image.

# RELU

1. Definition & Core Concept A ReLU activation is defined as

 $ReLU(x)=max(0, x).\text{ReLU}(x) = \max(0, x).ReLU(x)=max(0, x).$ 

It passes positive inputs unchanged and zeros out negatives, introducing non-linearity with a simple, piecewise-linear function.

2. Importance in AI/ML Pipelines

- Sparsity: Zeros out negative activations, producing sparse representations that can improve efficiency and generalization.
- Avoids Vanishing Gradients: For x>0x>0x>0, derivative is 1, so gradients don't shrink as they back-propagate through many layers.
  - Computational Simplicity: Very cheap to compute compared to sigmoid/tanh (no exponentials).

## 3. Key Variants & Techniques

- Leaky ReLU:  $max(\alpha x, x) max(\alpha x, x) max(\alpha x, x)$  with small  $\alpha = 0.01$  to keep a non-zero gradient for x < 0x < 0x < 0.01
  - Parametric ReLU (PReLU): Learns α\alphaα during training.
  - ELU (Exponential Linear Unit): Smooth negative side  $\alpha(ex-1)$ \alpha(e^x-1)\alpha(ex-1) to speed up learning.
- RReLU (Randomized Leaky ReLU): Uses a random slope for the negative side during training for regularization.

## 4. Advantages & Limitations

Advantages	Limitations
Simple, fast to compute	"Dead ReLU" problem: neurons stuck at 0
Mitigates vanishing-gradient via slope=1 for x>0x>0x>0	Not zero-centered: positive outputs only
Encourages sparse activations (efficiency)	Can suffer from high variance gradients

## 5. Real-World Applications & Examples

- Deep CNNs: Standard choice in architectures like ResNet, VGG, MobileNet.
- Feedforward Networks: Default for hidden layers in many MLPs.
- Generative Models: Used in GAN discriminators and many encoder/decoder blocks.

# 6. Sample Code Snippet (TensorFlow/Keras)

```
from tensorflow.keras import layers, models

model = models.Sequential([

layers.Dense(64, input_shape=(100,)),

layers.ReLU(), # explicit ReLU layer

layers.Dense(64),

layers.Activation('relu'), # alternative syntax

layers.Dense(10, activation='softmax')

])

model.compile(optimizer='adam',

loss='sparse_categorical_crossentropy',

metrics=['accuracy'])
```

## 7. Common Interview Questions & How to Answer

• Q: "Why choose ReLU over sigmoid?"

A: "ReLU avoids vanishing gradients for positive inputs and is computationally cheaper—leading to faster convergence in deep nets."

• Q: "What is the 'dying ReLU' problem?"

A: "If a ReLU neuron's weights update so it always outputs negative values, its gradient is zero forever—effectively killing that neuron."

• Q: "How do Leaky ReLU and PReLU address dying neurons?"

A: "They allow a small, non-zero gradient for x<0x<0x<0 (fixed  $\alpha$  in Leaky, learnable in PReLU) so neurons can recover."

• Q: "Are there cases where you wouldn't use ReLU?"

A: "In very shallow nets or when you need zero-centered outputs (tanh), or if negative activations carry meaning—then consider tanh or other activations."

## **Neural Networks Multiple Inputs and Outputs**

## 1. Definition & Core Concept

A multi-input, multi-output (MIMO) neural network accepts more than one distinct data stream and produces multiple targets in one forward pass. Internally, branches process each input modality (or feature group), then merge into a shared representation before possibly splitting again into separate heads for each output.

## 2. Importance in AI/ML Pipelines

- Unified Modeling: Learns interdependencies between heterogeneous data sources (e.g., images + text).
- Efficiency: Shares computation in early layers, reducing redundant training across separate models.
- Joint Optimization: Balances trade-offs between tasks (multi-task learning), improving generalization.

# 3. Key Architectures & Techniques

## 1. Branch-Merge-Head

- o Branches: Separate subnetworks per input (e.g., CNN for images, RNN for text).
- Merge: Concatenate or add branch outputs into a common dense block.
- Heads: Separate output layers for each target (e.g., one regression head, one classification head).

## 2. Shared-Backbone + Task-Specific Heads

A deep "backbone" encodes all inputs jointly; then fork into specialized heads for each output.

## 3. Cross-Stitch / Sluice Networks

Learn how much sharing vs. separation between tasks, via trainable "gating" between layers.

## 4. Multi-Modal Fusion

Early fusion (concatenate raw inputs), late fusion (combine high-level features), or hybrid.

## 4. Advantages & Limitations

Advantages	Limitations

Exploits complementary information across inputs	Requires careful design of branch/head sizes
Improves data efficiency via shared features	Loss balancing between outputs can be tricky
Reduces serving complexity vs. separate models	Risk of negative transfer (tasks hurting each other)

# 5. Real-World Applications & Examples

- Autonomous Driving: Fuse camera, lidar, and radar inputs; predict steering angle, speed, and object labels.
- Healthcare Diagnostics: Combine patient metadata (age, labs) + imaging (X-rays) to output disease probability and recommended treatments.
- Recommender Systems: Multi-modal inputs (user profile + item text + item image), outputs: click-through probability and purchase likelihood.
- Financial Forecasting: Use economic indicators + news sentiment; outputs: stock price prediction and risk score.

# 6. Sample Code Snippet (TensorFlow/Keras)

from tensorflow.keras import layers, models, Input

# # Define two inputs

image\_input = Input(shape=(224,224,3), name='img\_in')

text\_input = Input(shape=(100,), name='txt\_in')

## # Branch 1: CNN for images

x1 = layers.Conv2D(32,3,activation='relu')(image\_input)

x1 = layers.Flatten()(x1)

```
# Branch 2: Dense for text embeddings
x2 = layers.Embedding(input_dim=5000, output_dim=64)(text_input)
x2 = layers.GlobalAveragePooling1D()(x2)
# Merge
merged = layers.concatenate([x1, x2])
shared = layers.Dense(128, activation='relu')(merged)
# Head 1: classification
cls_out = layers.Dense(10, activation='softmax', name='class_out')(shared)
# Head 2: regression
reg_out = layers.Dense(1, activation='linear', name='value_out')(shared)
model = models.Model(inputs=[image_input, text_input],
           outputs=[cls_out, reg_out])
model.compile(optimizer='adam',
       loss={'class_out':'sparse_categorical_crossentropy',
           'value_out':'mse'},
       loss_weights={'class_out':1.0, 'value_out':0.5},
       metrics=['accuracy'])
```

# 7. Common Interview Questions & How to Answer

• Q: "How do you balance losses across outputs?"
A: "Use loss weights, dynamic weighting (uncertainty-based), or normalize losses by scale."

• Q: "When might you avoid a shared-backbone?"

A: "If tasks are too divergent-negative transfer may hurt performance; then use mostly separate branches."

Q: "How do you fuse very different modalities?"

A: "Choose early/late fusion based on data synchronization; often process each modality separately then merge high-level features."

• Q: "What's negative transfer and how to mitigate it?"

A: "When one task degrades another; mitigate via task-specific layers, gating (cross-stitch), or selective layer sharing."

### **Neural Networks ArgMax and SoftMax**

1. Definitions & Core Concepts

#### **Softmax**

- Purpose: Converts a raw score vector z∈RK\mathbf{z}\in\mathbb{R}^Kz∈RK into a probability distribution over KKK classes.
  - Formula:

 $softmax(zi) = zi\sum_{j=1}^{i=1,...,K} \operatorname{softmax}(z_i) = \frac{e^{z_i}}{\sum_{j=1}^{K} e^{z_j}} \operatorname{softmax}(z_i) = \frac{e^{z_i}}{\sum_{j=1}^{K} e^{z_j}} \operatorname{softmax}(z_i) = \lim_{j=1,...,K} \operatorname{softmax}(z_i) = \lim_$ 

- Key Properties:
  - Outputs sum to 1:  $\sum isoftmax(zi)=1\sum i \text{ } text{softmax}(z_i) = 1\sum isoftmax(zi)=1.$
  - Monotonic: larger ziz\_izi ⇒ larger probability.
  - o Smooth, differentiable—ideal for gradient-based learning.

#### **ArgMax**

- Purpose: Picks the index of the highest-scoring class. Often used at inference to choose the "winning" class.
- Definition:
   argmaxi zi={ i|zi≥zj ∀ j}. \arg\max\_i\,z\_i = \{\,i\mid z\_i \qe z\_i\;\forall\,j\}.argimaxzi={i | zi≥zj ∀ j}.
  - Key Point: ArgMax is non-differentiable, so it's only used after training (no gradients flow through it).

# 2. Why Softmax + ArgMax Together?

### 1. Training Phase:

- Use Softmax + Cross-Entropy Loss to teach the network to assign high probability to the correct class.
- Loss:
   L=-∑i=1Kyilog(softmax(zi)), L = -\sum\_{i=1}^K y\_i
   \log(\mathrm{softmax}(z\_i)), L=-i=1∑Kyilog(softmax(zi)),
   where yiy\_iyi is the one-hot label.

#### 2. Inference Phase:

- Compute raw logits z\mathbf{z}z.
- Apply Softmax to interpret as probabilities (e.g. for confidence scores).
- Use ArgMax on z\mathbf{z}z (or on softmax outputs—they give same ranking) to pick the predicted class

### 3. Numerical Stability Tricks

• Log-sum-exp trick: softmax(zi)=ezi-m $\sum$ jezj-m,m=maxjzj \mathrm{softmax}(z\_i) = \frac{e^{z\_i - m}}{\sum\_j e^{z\_j - m}}, \quad m = \max\_j z\_j softmax(zi)=\sum\_j ezj-mezi-m,m=jmaxzj prevents overflow when zzz has large values.

• Combined Softmax + Cross-Entropy APIs: Most frameworks provide a single "softmax-cross-entropy" op that fuses both steps for stability and performance.

### 4. Gradient of Softmax

For a single example with logits zzz and true class ttt, the derivative of the loss w.r.t. ziz\_izi is:

This simple form comes from combining the softmax Jacobian and cross-entropy derivative.

# 5. Sample Code Snippets

## PyTorch (training & inference)

```
import torch, torch.nn as nn, torch.nn.functional as F
logits = torch.randn(3, 5) # batch of 3, 5 classes
labels = torch.tensor([2,0,4]) # true class indices
# Training: softmax + cross-entropy
loss = F.cross_entropy(logits, labels)
# Inference: probabilities + prediction
probs = F.softmax(logits, dim=1) # shape (3,5), sums to 1
preds = torch.argmax(logits, dim=1) # shape (3,), class indices
TensorFlow/Keras
import tensorflow as tf
# model's final Dense layer typically omits activation:
logits = tf.random.normal((3,5))
# Training
loss = tf.reduce_mean(
  tf.nn.sparse_softmax_cross_entropy_with_logits(
    labels=[2,0,4], logits=logits
```

```
# Inference
probs = tf.nn.softmax(logits, axis=1) # probabilities
preds = tf.argmax(logits, axis=1) # class indices
```

#### 6. Common Interview Ouestions & Answers

• Q: "Why do we combine Softmax with Cross-Entropy, not just MSE?"

A: "Cross-entropy matches the probabilistic output of softmax and yields better gradients when probabilities are low/high; MSE on probabilities can lead to vanishing gradients."

• Q: "Why is ArgMax non-differentiable, and does it matter?"

A: "ArgMax's output jumps discreetly when the max index changes, so it has zero-gradient almost everywhere; but we only use it at inference, so it doesn't block training."

Q: "How does temperature scaling affect softmax?"

A: "Using softmax(zi/T)\mathrm{softmax}(z\_i/T)softmax(zi/T) with T>1T>1T>1 makes the distribution 'softer' (more uniform); T<1T<1 makes it peakier—useful in distillation or calibrating confidence."

• Q: "Can we backpropagate through ArgMax?"

A: "No—but you can approximate it with soft-argmax or use techniques like Gumbel-Softmax for differentiable sampling in structured models."

#### The SoftMax Derivative

1. Recap: Softmax

```
For a logits vector z=[z1,...,zK]\mathbf{z} = [z_1,\dots,z_K]z=[z1,...,zK],
```

 $\sigma(z)=ezi\sum_{j=1}Kezj$  for i=1,...,K. sigma\_i(\mathbf{z}) = \frac{e^{z\_i}}{\sum\_{j=1}Kez} \quad\text{for }i=1,\dots,K. $\sigma(z)=\sum_{j=1}Kez$  jezi for i=1,...,K.

### 2. Why Its Derivative Matters

• It's central to backprop when you combine Softmax with cross-entropy loss.

• You need the Jacobian ∂σiðzj\frac{\partial \sigma\_i}{\partial z\_j}∂zj∂σi to compute how a change in one logit affects all output probabilities.

### 3. Jacobian Matrix of Softmax

The derivative is a K×KK\times KK×K matrix:

 $\partial \sigma i \partial z j = \{\sigma i (1 - \sigma i), if i = j, \neg \sigma i \sigma j, if i \neq j, \frac \{\rho i = j, \ \} = \ \} = \ \{i = j, \ \} = \ \{i = j, \ \} = \ \} = \ \{i = j, \ \}$ 

Or in compact form:

 $\partial \sigma_i \partial z_j = \sigma_i(\delta_i - \sigma_j), \frac{(\delta_i - \sigma_j)}{\sigma_i} = \sigma_i(\delta_i - \sigma_j),$ 

where  $\delta ij \leq 1$  when i=j=j=j, else 0.

4. Combined Softmax + Cross-Entropy Simplification

For cross-entropy loss L= $-\Sigma$ iyilog( $\sigma$ i)L = -\sum\_{i} y\_i \log(\sigma\_i)L= $-\Sigma$ iyilog( $\sigma$ i), the gradient w.r.t. logits simplifies to:

 $\partial L \partial z i = \sigma i - \gamma i. \frac{\pi i}{\pi i} = \frac{i}{\pi i} = \sin \alpha_i - \gamma_i \partial z i \partial L = \sigma i - \gamma i.$ 

That is, softmax derivative and log-loss cancel out into a neat (p-yp - yp-y) term—no need to materialize the full Jacobian in code!

5. Worked Mini-Example (K=3)

Let z=[2.0,1.0,0.1]\mathbf{z} = [2.0, 1.0, 0.1]z=[2.0,1.0,0.1]. First compute softmax:

text

CopyEdit

 $e^{z} = [7.389, 2.718, 1.105] \rightarrow sum \approx 11.212$ 

 $\sigma = [0.659, 0.242, 0.099]$ 

Jacobian entries:

•  $\partial \sigma_1/\partial z_1 = 0.659 * (1 - 0.659) \approx 0.225$ 

```
• \partial \sigma_1/\partial z_2 = -0.659 * 0.242 \approx -0.160
```

```
• \partial \sigma_1/\partial z_3 = -0.659 * 0.099 \approx -0.065 ...and so on for rows 2,3.
```

```
6. Numerical Stability Tip
```

Always subtract max-logit before exponentiating:

```
z\sim i=zi-maxjzj.\tilde z_i=z_i-max_jz_j.z\sim i=zi-jmaxzj.
```

The Jacobian formula remains unchanged, but your  $\boldsymbol{\sigma}$  values stay in a safe numeric range.

```
7. Sample Code Snippet (NumPy)
```

python

CopyEdit

import numpy as np

```
def softmax(z):
```

```
z = z - np.max(z)
```

e = np.exp(z)

return e / np.sum(e)

def softmax\_jacobian(z):

```
s = softmax(z).reshape(-1,1) # column vector
```

return np.diagflat(s) - s @ s.T # diag(s) - s s^T

# Example

z = np.array([2.0,1.0,0.1])

J = softmax\_jacobian(z)

Drintt Jacobian:\n . J	print	("Jacobian:\n",	J)
------------------------	-------	-----------------	----

#### 8. Possible Interview Questions

• Q: "Derive the Jacobian of softmax."

A: Show you know the  $\sigma i(\delta ij - \sigma j) \simeq i(\delta ij - \sigma j) = i(\delta i$ 

• Q: "Why do we combine softmax+cross-entropy in one op?"

A: For numeric stability and computational efficiency—avoids underflow/overflow in log and exp.

• Q: "How would you backpropagate through a softmax layer manually?"

A: Multiply upstream gradient (vector) by the Jacobian matrix (or use the simplified p-yp - yp-y when using cross-entropy).

#### **Cross Entropy**

### 1. Definition & Core Concept

Cross-Entropy measures the disparity between two probability distributions: the true distribution ppp (often one-hot labels) and the model's predicted distribution qqq. For a single example over KKK classes:

 $H(p,q)=-\sum_{i=1}^{n}Kpi \log q_i.H(p,q)=-\sum_{i=1}^{n}Kpi \log q_i.H(p,q)=-i=1\sum_{i=1}^{n}Kpi \log q_i.H($ 

When ppp is one-hot (say class ttt), this simplifies to

 $H(p,q) = -\log qt.H(p,q) = -\log q_t.H(p,q) = -\log qt.$ 

#### 2. Why It Matters in AI/ML

- Training Objective: Preferred loss for classification tasks because it penalizes confident, wrong predictions heavily.
  - Probabilistic Interpretation: Equivalent to maximizing the log-likelihood of the correct class under the model.
- Gradient Simplicity: Combines neatly with softmax to yield  $\partial L \partial z = qi pi \frac{z_i}{q_i p_i \partial z_i} = q_i p_i \partial z_i \partial L = q_i q_i \partial z_i \partial L = q_i q_i \partial z_i \partial L = q_i \partial z_i \partial z_i \partial L = q_i \partial z_i \partial$

#### 3. Variants

Variant	Formula	Use Case
Binary Cross-Entropy	$ -[ylogy^++(1-y)log(1-y^*)]-[y\setminus y+(1-y)\setminus y+(1-$	Binary classification, sigmoid
Categorical Cross-Entropy	−Σi=1Kyilogy^i-\sum_{i=1}^K y_i \log \hat y_i−Σi=1Kyilogy^i	Multi-class, softmax output
Sparse Categorical CE	Same as categorical CE but labels as integer indices	Saves memory when KKK large
Weighted Cross-Entropy	−∑iwi yilogy^i-\sum_i w_iy_i\log \hat y_i−∑iwiyilogy^i	Imbalanced classes

# 4. Mathematical Properties

- Non-negative:  $H(p,q) \ge 0H(p,q) \setminus geoH(p,q) \ge 0$ , with equality iff p=qp=qp=q exactly.
- Asymmetric:  $H(p,q) \neq H(q,p)H(p,q) \setminus H(q,p)H(p,q) = H(q,p)$ ; measures how well qqq approximates ppp.
- Convex in qqq: Guarantees a single global minimum when optimizing.

# 5. Sample Code Snippets

import torch, torch.nn.functional as F

logits = torch.randn(4, 3) # batch of 4, 3 classes

labels = torch.tensor([0,2,1,2])

# Categorical cross-entropy

loss = F.cross\_entropy(logits, labels)

```
TensorFlow/Keras
python
CopyEdit
import tensorflow as tf
# logits shape (batch, classes), labels integer indices
loss = tf.reduce_mean(
 tf.nn.sparse_softmax_cross_entropy_with_logits(
    labels=[0,2,1,2], logits=logits
  )
NumPy (manual softmax + CE)
python
CopyEdit
import numpy as np
def softmax(z):
  e = np.exp(z - np.max(z))
 return e / e.sum()
def cross_entropy(p_true, p_pred):
 return -np.sum(p_true * np.log(p_pred))
# Example
y_{true} = np.array([0,1,0])
                             # one-hot for class 2
z = np.array([2.0, 1.0, 0.1])
y_pred = softmax(z)
print("Loss:", cross\_entropy(y\_true, y\_pred))
```

- 6. Common Interview Questions & How to Answer
- Q: "Why use cross-entropy instead of mean squared error for classification?" A: "Cross-entropy aligns with the probabilistic interpretation of classification, penalizes confident errors more sharply, and provides stronger, smoother gradients."
- Q: "How does binary cross-entropy relate to logistic regression?" A: "Binary CE is exactly the negative log-likelihood for a Bernoulli logistic model."
  - Q: "What's the advantage of sparse categorical CE?"

A: "Stores labels as integers instead of one-hot vectors-more memory- and compute-efficient when there are many classes."

• Q: "How would you handle class imbalance?"

A: "Use weighted cross-entropy or focal loss to up-weight rare classes."

## **Cross Entropy Derivatives and Backpropagation**

- 1. Recap of Cross-Entropy Loss
- Binary CE (for a single example):  $L = -[y\log(y^*) + (1-y)\log(1-y^*)], L = -\frac{y\log(y^*) + (1-y)\log(1-y^*)}{L} = -\frac{y\log(y^*) + (1-y)\log(y^*)}{L} = -\frac{y\log(y^*) + (1-y)\log(y^*)}{L} = -\frac{y\log(y^*) + (1-y)\log(y^*)}{L} = -\frac{y\log(y^*)}{L} = -\frac{y\log($ where  $y^* = \sigma(z) \cdot y = sigma(z) \cdot y^* = \sigma(z)$ .
- Categorical CE (one-hot labels, logits ziz\_izi):  $L = -\sum_{i=1}^{n} Ky_i \log(softmax(z)_i). L = -\sum_{i=1}^{n} Ky_i \log(softmax(z)_i). L = -i = 1\sum_{i=1}^{n} Ky_i \log(s$
- 2. Derivative of Binary Cross-Entropy

Given zzz (logit),  $y^=\sigma(z)$  (hat  $y = \sigma(z)$ , and true label  $y \in \{0,1\}y \in \{0,1\}y \in \{0,1\}$ ):

- 1. Derivative w.r.t. prediction  $\partial L \partial y^{--yy^{+1-y1-y^{-}}} - y^{-y^{+1-y1-y^{-}}} = -\frac{1-y}{1-\frac{1-y}{1-\frac{1-y}{1-\frac{1-y}{1-y^{-}}}}}$
- 2. Link to logit via sigmoid'  $dy^dz=y^(1-y^)\cdot dfrac\{d\cdot y\}\{dz\} = \cdot y\cdot (1-\cdot y^1-\cdot y^2).$

3.	Chain rule				
∂L∂z=∂	$L\partial y^{dy^{dz}=(y^{-}y).$	\frac{\partial L}{\partial z}	= \frac{\partial L}{\partial	$\hat{y}_{\pi}$	= \bigl(\ha
y - y\big	r).∂z∂L=∂y^∂Ldzdy	^=(y^-y).			

Key insight: Binary CE & sigmoid combine to yield y^-y\hat y - yy^-y.

## 3. Derivative of Categorical Cross-Entropy

For logits zjz\_jzj, softmax outputs pj=softmax(z)jp\_j = \mathrm{softmax}(z)\_jpj=softmax(z)j, and one-hot label yjy\_jyj:

- 1. Jacobian of softmax  $\partial pi\partial z_j = p_i(\delta i_j p_j). \partial z_j \partial p_i = p_i(\delta i_j p_j). \partial z_j \partial p_i = p_i(\delta i_j p_j). \partial z_j \partial p_i = p_i(\delta i_j p_j).$
- 2. Loss derivative w.r.t. softmax output  $\partial L\partial pi=-yipi.\displaystyle \frac{p_i}{p_i}.\partial pi\partial L=-piyi.$
- 3. Chain rule  $\rightarrow$  logits  $\partial L\partial z_j = \sum_{i=1}^{L} \Delta z_j = \sum_{i=1}^{K} \frac{1}{\pi c_{\beta_i}} = \sum_{i=1}^$

Key insight: Categorical CE + softmax collapse to p-yp - yp-y per class.

### 4. Backpropagation Integration

In a network, you typically have:

- Logits layer outputs z\mathbf{z}z.
- Loss node computes L(z)L(\mathbf{z})L(z).

During backprop, you pass upstream gradient  $\partial L/\partial z$  (vector of pj-yjp\_j - y\_jpj-yj) into the previous Dense layer:

 $\nabla$ WL=( $\partial$ L/ $\partial$ z) aprevT, $\nabla$ bL= $\partial$ L/ $\partial$ z,\nabla\_W L = (\partial L/\partial z)\; a\_{\text{prev}}^T,\quad \nabla\_b L = \partial L/\partial z, $\nabla$ WL=( $\partial$ L/ $\partial$ z)aprevT, $\nabla$ bL= $\partial$ L/ $\partial$ z,

where apreva\_{\text{prev}}aprev are activations from the layer below.

# 5. Worked Mini-Example (Binary)

- Logit z=0.2z=0.2z=0.2, label y=1y=1y=1.
- $y^{-\sigma}(0.2)\approx 0.55$ \hat y = \sigma(0.2)\approx0.55 $y^{-\sigma}(0.2)\approx 0.55$ .
- $\partial L/\partial z=y^-y=-0.45$ .\partial L/\partial z = \hat y y = -0.45. $\partial L/\partial z=y^-y=-0.45$ .

If this logit came from weight www and input x=0.7x=0.7x=0.7, then

6. Sample Code Snippet (PyTorch)

python

CopyEdit

import torch, torch.nn.functional as F

# logits and labels

logits = torch.tensor([[0.2, -1.0], [1.5, 0.3]], requires\_grad=True)

labels = torch.tensor([0, 1]) # two samples

# loss

loss = F.cross\_entropy(logits, labels)

loss.backward()

print("Gradients w.r.t. logits:\n", logits.grad)

# yields tensor([[p0 - y0, p1 - y0],

# [p0 - y1, p1 - y1]])

## 7. Common Interview Questions & How to Answer

• Q: "Why does CE + activation simplify gradients to p-yp - yp-y?"

A: "Because the derivative of the log cancels the softmax's denominator, and the chain rule folds neatly into pi-yip\_i - y\_ipi-yi."

• Q: "How does this help training efficiency?"

A: "You avoid computing full Jacobians—just a simple subtraction per class, which is fast and numerically stable."

• Q: "What about when using label smoothing?"

A: "You replace yyy with a softened distribution (e.g.  $yi=0.9y_$ 

• Q: "How do you implement weighted CE for imbalanced data?"

A: "Multiply each term -yilogpi-y\_i\log p\_i-yilogpi by a class weight wiw\_iwi; gradient becomes wi (pi-yi)w\_i\,(p\_i - y\_i)wi(pi-yi)."

#### Image Classification with Convolutional Neural Networks (CNNs)

### 1. Definition & Core Concept

A Convolutional Neural Network is a specialized feedforward network designed to process grid-structured data (like images). It uses learnable convolutional filters that slide over the input to detect local patterns, then builds hierarchical representations through stacked layers.

### 2. Why It Matters in AI/ML Pipelines

- Automatic Feature Extraction: Learns edge, texture, shape, and object detectors without manual engineering.
- Parameter Sharing: Convolutional kernels are reused spatially—dramatically reduces parameters vs. dense nets.
- Translation Invariance: Pooling and local receptive fields help the model recognize objects regardless of position.

# 3. Key Components & Techniques

- 1. Convolutional Layers
  - Apply k × times k×k kernels over input channels to produce feature maps.

Learn filters that respond to edges, colors, and complex textures.

#### 2. Activation Functions

• ReLU (or variants) applied after convolutions for non-linearity.

### 3. Pooling Layers

Pooling layer is used in CNNs to reduce the spatial dimensions (width and height) of the input feature maps while retaining the most important information.

- Max-Pooling (common): downsamples by taking local maximums.
- o Average-Pooling: smooths outputs—sometimes used before classification.

#### 4. Batch Normalization

o Normalizes activations each mini-batch to accelerate convergence and improve stability.

# 5. Dropout

o Randomly zeros activations to reduce overfitting in fully-connected or conv layers.

# 6. Fully-Connected (Dense) Head

Flattens feature maps and maps to class scores via one or more dense layers.

#### 4. Popular Architectures

- LeNet-5 (1998): Pioneering small CNN for digit recognition.
- AlexNet (2012): Eight-layer net that won ImageNet—popularized ReLU, dropout, data augmentation.
- VGG (2014): Deep stacks of 3×33\times33×3 conv layers—demonstrated depth matters.
- ResNet (2015): Introduced residual ("skip") connections to train very deep nets (50–152 layers).

• EfficientNet (2019): Compound scaling of depth, width, and resolution for optimal efficiency.

## 5. Advantages & Limitations

Advantages	Limitations
Learns rich, hierarchical image features	Requires large labeled datasets (e.g., ImageNet)
Parameter-efficient via weight sharing	High compute cost—especially deeper models
Strong translation invariance	Poor at modeling long-range dependencies without design tweaks
State-of-the-art on vision benchmarks	Often seen as a "black box"—interpretability issues

## 6. Real-World Applications & Examples

- ImageNet Classification: Benchmarks hundreds of object categories (e.g., dogs, cars) √
- Medical Imaging: Detects tumors in MRI/CT scans with high sensitivity.
- Autonomous Vehicles: Classify traffic signs, pedestrians, and drivable areas.
- Retail & Surveillance: Product recognition, face detection, anomaly spotting.
- Agriculture: Crop/disease identification from aerial or field images.

# 7. Sample Code Snippet (TensorFlow/Keras)

from tensorflow.keras import layers, models

def build\_simple\_cnn(input\_shape, num\_classes):

model = models.Sequential([

```
layers.Input(shape=input_shape),
    layers.Conv2D(32, (3,3), activation='relu'),
    layers.MaxPooling2D((2,2)),
    layers.Conv2D(64, (3,3), activation='relu'),
    layers.BatchNormalization(),
    layers.MaxPooling2D((2,2)),
    layers.Conv2D(128, (3,3), activation='relu'),
    layers.Flatten(),
    layers.Dropout(0.5),
    layers.Dense(128, activation='relu'),
    layers.Dense(num_classes, activation='softmax')
  ])
  model.compile(optimizer='adam',
         loss='sparse_categorical_crossentropy',
         metrics=['accuracy'])
  return model
# Example usage:
# model = build_simple_cnn((64,64,3), 10)
# model.fit(train_images, train_labels, epochs=20, batch_size=32, validation_split=0.1)
```

# 8. Common Interview Questions & How to Answer

• Q: "Why use pooling layers?"

A: "To reduce spatial dimensions, control overfitting, and build invariance to small translations."

• Q: "How do residual connections help?"

A: "By allowing gradients to skip layers, they alleviate vanishing gradients and enable very deep networks."

• Q: "When would you choose depthwise separable convolutions?"

A: "For mobile/embedded use cases (e.g., MobileNet), they reduce computation by factorizing spatial and channel convolutions."

• Q: "How do you prevent overfitting in CNNs?"

A: "Use data augmentation, dropout, weight decay (L2), early stopping, and transfer learning from pre-trained models."

#### **Residuals and Residual Blocks**

Residual (Statistics) = The difference between the actual value and the predicted value.

- yi: actual target value
- y^i: predicted value by your model

Residuals are most important in regression tasks, such as:

- Linear regression
- Polynomial regression
- Any supervised model outputting continuous values

Residual Value	Meaning
0	Perfect prediction
> 0	Underprediction (model too low)
< 0	Overprediction (model too high)

## **Residual Blocks (Residuals in Resnet)**

In ResNet, the term "residual" refers to residual blocks — a completely different concept than statistical residuals.

Problem it solves:

Deep networks suffer from vanishing gradients and degradation (more layers = worse accuracy)

Solution:

Instead of learning a full function F(x), let the network learn only the residual F(x)=H(x)-x

Then:

$$H(x)=F(x)+x$$

#### Where:

- x: input
- F(x): output of the residual block (usually 2 or 3 conv layers)
- H(x): final output with skip connection

#### Benefits:

- Easier optimization
- Enables very deep networks (ResNet-50, ResNet-101)
- Encourages identity mapping if learning unnecessary layers

# GANs

#### 1. Definition & Core Concept

A Generative Adversarial Network consists of two neural nets—the Generator (G) and the Discriminator (D)—that play a minimax game:

- G learns to map random noise  $z\sim pzz \sim pz$  to realistic samples G(z)G(z)G(z).
- D learns to distinguish real data x-pdatax\sim p\_{\text{data}}x-pdata from fake samples G(z)G(z)G(z). The objective is

 $\label{logD} \mbox{$minGmaxD Ex-pdata[logD(x)]+Ez-pz[log(1-D(G(z)))].\mbox{$min_G \max_D \ \mbox{$max_D \ \nox{$max_D \ \nox{$ 

2. Why	Ιt	Matters	in	AI/	'ML
--------	----	---------	----	-----	-----

- Powerful Generative Modeling: Can produce highly realistic images, audio, and more.
- Unsupervised Learning: Learns data distribution without labels.
- Foundation for Advances: Underlies StyleGAN (face synthesis), CycleGAN (unpaired translation), and beyond.

## 3. Key Architectures & Variants

- 1. DCGAN (Deep Convolutional GAN)
  - o Uses strided conv/deconv and batch-norm—stabilized training for images.
- 2. Conditional GAN (cGAN)
  - $\circ$  Conditions G and D on labels yyy: G(z,y)G(z,y)G(z,y), D(x,y)D(x,y)D(x,y). Enables class-conditional generation.
- 3. Wasserstein GAN (WGAN & WGAN-GP)
  - Replaces JS-divergence with Earth-Mover distance; uses weight-clipping or gradient penalty for stability.

# 4. CycleGAN

- o Learns unpaired image-to-image translation via cycle-consistency losses.
- 5. StyleGAN / StyleGAN2
  - o Introduces style-based generator with adaptive instance norm for fine control over synthesis.

# 4. Advantages & Limitations

Advantages	Limitations
Produces sharp, high-fidelity samples	Training can be unstable—mode collapse, oscillation
Unsupervised—no paired labels needed (many variants)	Sensitive to hyperparameters and architecture tweaks
Flexible: easy to condition, chain, or adapt	Evaluation metrics (IS, FID) imperfect

# 5. Real-World Applications & Examples

- Image Synthesis: Photorealistic faces (StyleGAN), artwork creation.
- Data Augmentation: Generate rare classes for medical imaging.
- Domain Translation: Photos↔Paintings (CycleGAN), day↔night scenes.
- Super-Resolution: Enhance image resolution (SRGAN).
- Anomaly Detection: Train on "normal" data—GAN's inability to reproduce anomalies flags outliers.

# 6. Sample Code Snippet (PyTorch - Simple DCGAN Generator & Discriminator)

import torch

import torch.nn as nn

# Generator

class Generator(nn.Module):

```
def __init__(self, nz, ngf, nc):
    super().__init__()
    self.main = nn.Sequential(
      nn.ConvTranspose2d(nz, ngf*8, 4, 1, 0, bias=False),
      nn.BatchNorm2d(ngf*8), nn.ReLU(True),
      nn.ConvTranspose2d(ngf*8, ngf*4, 4, 2, 1, bias=False),
      nn.BatchNorm2d(ngf*4), nn.ReLU(True),
      nn.ConvTranspose2d(ngf*4, ngf*2, 4, 2, 1, bias=False),
      nn.BatchNorm2d(ngf*2), nn.ReLU(True),
      nn.ConvTranspose2d(ngf*2, ngf, 4, 2, 1, bias=False),
      nn.BatchNorm2d(ngf), nn.ReLU(True),
      nn.ConvTranspose2d(ngf, nc, 4, 2, 1, bias=False),
      nn.Tanh()
    )
  def forward(self, input):
    return self.main(input)
# Discriminator
class Discriminator(nn.Module):
  def __init__(self, nc, ndf):
    super().__init__()
    self.main = nn.Sequential(
      nn.Conv2d(nc, ndf, 4, 2, 1, bias=False),
      nn.LeakyReLU(0.2, inplace=True),
      nn.Conv2d(ndf, ndf*2, 4, 2, 1, bias=False),
      nn.BatchNorm2d(ndf*2), nn.LeakyReLU(0.2, inplace=True),
      nn.Conv2d(ndf*2, ndf*4, 4, 2, 1, bias=False),
```

```
nn.BatchNorm2d(ndf*4), nn.LeakyReLU(0.2, inplace=True),
nn.Conv2d(ndf*4, ndf*8, 4, 2, 1, bias=False),
nn.BatchNorm2d(ndf*8), nn.LeakyReLU(0.2, inplace=True),
nn.Conv2d(ndf*8, 1, 4, 1, 0, bias=False),
nn.Sigmoid()
)

def forward(self, input):
return self.main(input).view(-1, 1).squeeze(1)
```

#### 7. Training Loop Sketch

- 1. Update D: maximize  $\log D(x) + \log (1 D(G(z))) \log D(x) + \log (1 D(G(z))) \log D(x) + \log (1 D(G(z)))$
- 2. Update G: minimize log(1-D(G(z))) log(1-D(G(z))) log(1-D(G(z))) (or maximize <math>logD(G(z)) logD(G(z)) logD(
  - 3. Alternate G/D steps, use label smoothing, batch-norm, learning-rate schedules.
- 8. Common Interview Questions & How to Answer
  - Q: "Why is GAN training unstable?"

A: "Because G and D objectives clash—if one overtakes, the other's gradients vanish; balancing capacity, learning rates, and using WGAN or spectral norm helps."

• Q: "What is mode collapse and how to mitigate it?"

A: "G maps diverse zzz to the same output—mitigated via minibatch discrimination, historical averaging, or unrolled GANs."

• Q: "How does WGAN improve upon vanilla GAN?"

A: "Uses Wasserstein distance for smoother gradient signal; enforces 1-Lipschitz via weight clipping or gradient penalty."

• Q: "Explain conditional GANs."

A: "Inject labels into G and D (e.g., as extra channels or embedding) so generation is class-controlled."

•	Q: "How do you evaluate GAN quality?"
A: "Use	e Inception Score (IS). Frechet Inception Distance (FID), and human evaluation—each has trade-offs.

## **Tensors for Neural Networks**

₩ What Is a Tensor?

A tensor is a multi-dimensional array — a generalization of scalars, vectors, and matrices.

Tensor Type	Dimensions	Example	Shape
Scalar	0D	5	0
Vector	1D	[1, 2, 3]	(3,)
Matrix	2D	[[1, 2], [3, 4]]	(2, 2)
3D Tensor	3D	RGB image = [H, W, C]	(256, 256, 3)
nD Tensor	nD	Used in deep learning (e.g., batch of images)	(batch, H, W, C) or (batch, features)

A tensor is the input, output, and weight carrier throughout your entire neural network.

Why Tensors in Neural Networks?

Neural networks are layers of mathematical operations (like dot products, convolutions, etc.) applied to structured numerical data.

#### Tensors allow:

- Efficient representation of multi-dimensional data (images, audio, video, etc.)
- Parallel computation on GPUs using frameworks like TensorFlow or PyTorch
- Clean abstraction of complex data flow
- Tensor Shapes in Practice
- Dense (Fully Connected) Layer

## Input:

Shape: (batch\_size, features) → e.g., (64, 128)

# Weights:

• Shape: (128, 256) → transforms 128-dim into 256-dim

# Output:

- Shape: (64, 256)
- Image Tensor (For CNNs)
  - Single RGB image: (Height, Width, Channels)  $\rightarrow$  e.g., (224, 224, 3)
  - Batch of images: (Batch, Height, Width, Channels) → e.g., (32, 224, 224, 3)

PyTorch uses channel-first (B, C, H, W); TensorFlow uses channel-last (B, H, W, C)

\* Tensors Undergo Operations

In forward and backward passes, tensors go through:

- 1. Linear Transformation: Dot products (W\*x + b)
- 2. Non-Linearity: Apply activation function like ReLU(x)
- 3. Loss Calculation: Compare output tensor to labels
- 4. Backpropagation: Compute gradient tensors
- 5. Parameter Update: Use gradient tensors to update weight tensors

# Broadcasting: Tensor Magic

Broadcasting lets you apply operations between tensors of different shapes, e.g.:

$$[3, 4] + [1] \rightarrow [4, 5]$$

This helps neural networks efficiently manipulate entire batches of data.

# PyTorch Example

import torch

# NumPy vs Tensor Libraries

- ullet NumPy o Good for CPU, no autodiff
- TensorFlow, PyTorch → GPU support, automatic differentiation, training-ready

# Interview Talking Points

Q: What is a tensor in deep learning?

A: A general-purpose n-dimensional array used to represent data and model parameters.

• Q: Why are tensors important?

A: They enable parallel computation on GPUs and represent complex, structured data.

• Q: What's the difference between a tensor and a matrix?

A: A matrix is 2D; a tensor is n-dimensional (can represent matrices, vectors, scalars, etc.).

• Q: How are tensors stored and computed?

A: They're stored in contiguous memory blocks and optimized for hardware like GPUs and TPUs.

### **Vanishing and Exploding Gradients**

They refer to what happens to gradients during backpropagation in deep neural networks:

- $\bullet \quad \text{Vanishing gradients} \rightarrow \text{gradients become very small} \rightarrow \text{early layers stop learning}$
- $\bullet \quad \textbf{Exploding gradients} \rightarrow \textbf{gradients} \ become \ \textbf{very large} \rightarrow \textbf{unstable training} \ (\textbf{loss goes to NaN})$

These issues are especially problematic in:

- Very deep networks
- Recurrent Neural Networks (RNNs)

Why Does This Happen?

- 1. Activation Functions
  - Sigmoid / Tanh squish inputs to a small range → gradients are < 1 → repeated multiplication causes vanishing
  - ReLU helps with this because its gradient is either 0 or 1 (but it has its own issue "dying ReLUs")
- 2. Poor Weight Initialization
  - If weights are too small  $\rightarrow$  gradient shrinkage If too large  $\rightarrow$  gradient explosion
- 3. Deep Architectures

ullet The more layers you have, the more often you multiply gradients ullet exponential effects

## **X** Effects

Problem	Symptoms
Vanishing Gradients	Network stops learning, loss plateaus, especially early layers
Exploding Gradients	Model diverges, loss becomes NaN, weights blow up

# **%** Solutions

- 1. Use Better Activations
  - Replace sigmoid/tanh with ReLU, LeakyReLU, GELU
- 2. Use Proper Initialization
  - Xavier/Glorot Initialization: good for tanh
  - He Initialization: ideal for ReLU
- 3. Batch Normalization
  - Normalizes layer inputs → keeps gradients stable
- 4. Gradient Clipping (esp. in RNNs)
  - Clip gradients to a max value to prevent explosion:

torch.nn.utils.clip\_grad\_norm\_(model.parameters(), max\_norm=1.0)

- 5. Residual Connections (ResNet)
  - Let gradients bypass some layers → easier flow backward

#### **Recurrent Neural Networks (RNNs)**

1. Definition & Core Concept

A Recurrent Neural Network (RNN) is a type of neural network specially designed to handle sequences. Unlike feedforward networks, RNNs retain a memory of past inputs by using a hidden state that is updated at each timestep.

Given a sequence of inputs  $x1,x2,...,xTx_1,x_2$ , \dots,  $x_Tx_1,x_2,...,xT$ , the RNN updates as follows:

$$h_t = f(W_{xh}x_t + W_{hh}h_{t-1} + b_h)$$

- Xt: input at time step t
- ht: hidden state at time t
- h(t-1): hidden state from the previous time step
- Wxh: input-to-hidden weight matrix
- Whh: hidden-to-hidden (recurrent) weight matrix
- bh: bias
- f: nonlinearity (usually tanh or ReLU)

### **Key Features of RNNs:**

Sequential Data Processing:

RNNs are specifically designed to handle data that has a **sequential** structure, such as text, speech, or time series data.

## Hidden State:

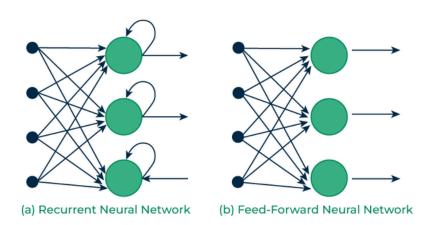
**RNNs have a hidden state that stores information from previous inputs**, allowing them to process sequential data effectively.

## Feedback Loops:

RNNs use feedback loops, which enable them to maintain context across different time steps in a sequence.

# **Shared Weights:**

**RNNs share weights** across different time steps, making them more efficient and capable of learning long-range dependencies.



# 2. Importance in AI/ML Pipelines

- Sequence Modeling: Natural for text, time series, audio, and any ordered data.
- Temporal Dependencies: Captures relationships across time (e.g., language context, trend patterns).
- Foundation for Advanced Models: Lays groundwork for LSTM, GRU, and attention-based Transformers.

#### 3. Key Architectures & Techniques

## 1. Vanilla RNN

Uses simple activation (tanh or ReLU) for recurrence; prone to vanishing/exploding gradients.

# 2. Long Short-Term Memory (LSTM)

Introduces gates (input, forget, output) and cell state to learn long-range dependencies.

## 3. Gated Recurrent Unit (GRU)

Simplifies LSTM with only reset and update gates—fewer parameters, faster training.

#### 4. Bidirectional RNN

Processes sequence forward and backward, then concatenates hidden states for richer context.

### 5. Stacked (Deep) RNNs

Multiple recurrent layers stacked to learn hierarchical temporal features.

### 6. Sequence-to-Sequence (Seq2Seq)

• Encoder–Decoder setup for tasks like translation, summarization, with optional attention.

# 4. Advantages & Limitations

Advantages	Limitations
Naturally handles variable-length sequences	Struggles with very long dependencies without gates
Parameter sharing across timesteps	Training can be slow due to sequential nature
Powerful for time-series & language modeling	Vanishing/exploding gradients in vanilla RNNs
Bidirectional & gated variants mitigate issues	Harder to parallelize than CNNs/Transformers

- 5. Real-World Applications & Examples
  - Language Modeling & Text Generation: Predict next word/character in a corpus.
  - Machine Translation: Seq2Seq with attention for translating sentences.
  - Speech Recognition: Convert audio frames to phonemes or text.
  - Time-Series Forecasting: Stock prices, weather, sensor readings.
  - Video Analysis: Action recognition by modeling frame sequences.

#### 6. Sample Code Snippet (Simple RNN Cell)

import numpy as np

```
def rnn_step(x_t, h_prev, W_xh, W_hh, b_h):
    return np.tanh(np.dot(W_xh, x_t) + np.dot(W_hh, h_prev) + b_h)
# Example usage:
# model = build_lstm_model((100, 50), 10)
```

# model.fit(x\_train, y\_train, epochs=20, batch\_size=32, validation\_split=0.1)

#### 7. Common Interview Questions & How to Answer

• Q: "Why do vanilla RNNs struggle with long-range dependencies?"

A: "Repeated multiplications of gradients (<1 or >1) across many timesteps cause vanishing or exploding gradients."

• Q: "When do you use a GRU over an LSTM?"

A: "When you need fewer parameters and faster training but still want most gating benefits."

Q: "What's the benefit of a bidirectional RNN?"

A: "It uses future context as well as past, improving predictions especially in text where both sides matter."

• Q: "How does attention improve Seg2Seg models?"

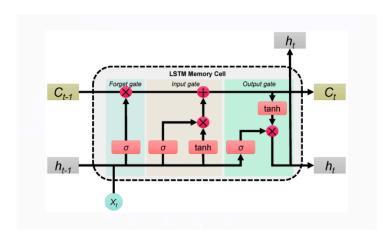
A: "It lets the decoder selectively focus on relevant encoder hidden states, alleviating information bottleneck."

# Long Short-Term Memory (LSTM)

## 1. Definition & Core Concept

An LSTM is a gated RNN variant designed to **remember long-term dependencies** and solve the **vanishing gradient problem** that standard RNNs suffer from.

It does this by introducing a **cell state** and a set of **gates** that control the flow of information.



# **LSTM Cell Overview:**

At each timestep t, the LSTM maintains:

- Cell state ct: long-term memory
- **Hidden state** ht: short-term memory / output
- Input xt

# Four key components (gates):

$$\begin{split} f_t &= \sigma(W_f \cdot [h_{t-1}, x_t] + b_f) & \text{Forget gate} \\ i_t &= \sigma(W_i \cdot [h_{t-1}, x_t] + b_i) & \text{Input gate} \\ \tilde{c}_t &= \tanh(W_c \cdot [h_{t-1}, x_t] + b_c) & \text{Candidate cell state} \\ c_t &= f_t \odot c_{t-1} + i_t \odot \tilde{c}_t & \text{New cell state} \\ o_t &= \sigma(W_o \cdot [h_{t-1}, x_t] + b_o) & \text{Output gate} \\ h_t &= o_t \odot \tanh(c_t) & \text{New hidden state} \end{split}$$

#### Where:

- σ: sigmoid function
- tanh: tanh activation function
- o: element-wise multiplication

## Intuition behind gates:

Gate	Purpose
Forget gate ft	Decides what to discard from cell state
Input gate it	Decides what new info to add
Candidate ct	Proposes <b>new values</b> to add to memory
Output gate ot	Decides what part of memory to output

## 2. Why It Matters in AI/ML

- Long-Range Memory: Gates mitigate vanishing/exploding gradients, letting the network learn dependencies across many timesteps.
  - Versatility: Core building block for sequence tasks where context over dozens or hundreds of steps matters.
  - Foundation for Attention: Often used as the encoder/decoder in attention-based Seq2Seq models.

## 3. Key Variants & Techniques

- Bidirectional LSTM: Processes sequence both forward and backward; concatenates or sums the two hidden states.
  - Stacked LSTM: Multiple LSTM layers layered for hierarchical temporal feature extraction.
  - Peephole Connections: Gates also see the previous cell state ct-1c\_{t-1}ct-1, allowing more precise timing.
  - Coupled Forget/Input Gate: Merges ftf\_tft and iti\_tit into a single gate to reduce parameters.
  - Layer Normalization in LSTM: Applies layer-norm to gates for faster convergence and stability.

### 4. Advantages & Limitations

Advantages	Limitations
Learns very long dependencies via gated pathways	More parameters and compute than vanilla RNNs
Resilient to vanishing/exploding gradients	Sequential by nature—slower to train than transformers on GPUs
Proven performance on many sequence tasks	Harder to parallelize across timesteps
Flexible: works in encoder-decoder and attention models	Can overfit small datasets without regularization

## 5. Real-World Applications & Examples

- Machine Translation: Seq2Seq LSTM with attention (e.g., early Google Translate models).
- Speech Recognition: Maps audio frames to text sequences (DeepSpeech 1 used LSTMs).
- Time-Series Forecasting: Predicting stock prices, weather, energy consumption.
- Text Generation: Language models that compose poetry or code (pre-Transformer era).
- Video Captioning: Encode video frame features over time, then decode to language.

## 6. Sample Code Snippet (TensorFlow/Keras)

import torch

import torch.nn as nn

lstm = nn.LSTM(input\_size=10, hidden\_size=20, num\_layers=1, batch\_first=True)

x = torch.randn(5, 3, 10) # (batch, seq\_len, input\_size)

output,  $(h_n, c_n) = Istm(x)$ 

#### 7. Common Interview Questions & How to Answer

Q: "Explain the role of each LSTM gate."

A: "Forget gate discards irrelevant past memory; input gate controls new information addition; output gate determines what part of cell state is exposed as hidden state."

Q: "Why use LSTM over vanilla RNN?"

A: "Its gated cell state avoids vanishing/exploding gradients, enabling learning of long-term dependencies."

• Q: "How do you initialize LSTM weights?"

A: "Use orthogonal initialization for recurrent kernels, Glorot or He for input kernels, and biases set so forget gate bias is often initialized to 1."

Q: "When might you choose a GRU instead?"

A: "When you need a simpler, faster-to-train model with fewer parameters but still want many gating benefits."

Q: "How do you prevent overfitting in LSTMs?"

A: "Apply dropout (both on inputs and recurrent connections), L2 weight decay, and consider reducing hidden size or early stopping."

#### NLP

#### 1. Definition & Core Concept

Natural Language Processing is the field at the intersection of computer science, linguistics, and AI that enables machines to understand, interpret, and generate human language. It transforms raw text or speech into structured representations (tokens, embeddings) that models can work with.

#### 2. Importance in AI/ML Pipelines

- Human-Machine Interaction: Powers chatbots, virtual assistants, and voice interfaces.
- Unstructured Data: Converts text—emails, social media, documents—into features for analytics and decision-making.

• Foundation for Language Models: Underlies everything from simple bag-of-words classifiers to large pre-trained Transformers.

### 3. Key Techniques & Components

## 1. Tokenization & Text Preprocessing

- o Splitting text into words/subwords (e.g., WordPiece, Byte-Pair Encoding).
- o Normalization (lowercasing, punctuation removal, stemming/lemmatization).

## 2. Feature Representations

- o Count-based: Bag-of-Words, TF-IDF.
- Dense Embeddings: Word2Vec, GloVe, FastText.
- o Contextual Embeddings: BERT, GPT, RoBERTa (deep contextualized vectors).

### 3. Core Architectures

- RNNs/LSTMs/GRUs: Early sequence models.
- o CNNs for NLP: Text classification via convolutions over word embeddings.
- o Transformers: Self-attention blocks that capture global context efficiently.

### 4. Common Tasks

- o Classification: Sentiment analysis, spam detection.
- Sequence Tagging: Named-entity recognition, POS tagging.
- o Generation: Machine translation, summarization, dialogue systems.
- o Retrieval & QA: Information retrieval, question answering over documents.

# 4. Advantages & Limitations

Advantages	Limitations
Handles diverse text data at scale	Requires large corpora and compute for pre-training
Transferable via fine-tuning	May inherit biases present in training data

Supports end-to-end pipelines	Interpretability remains challenging
Strong benchmarks across many languages	Struggles with low-resource or code-mixed text

## 5. Real-World Applications & Examples

- Chatbots & Virtual Assistants: Alexa, Siri, Google Assistant.
- Customer Analytics: Automated ticket triage, sentiment dashboards.
- Healthcare: Clinical note understanding, medical coding.
- Search & Recommend: Query expansion, document ranking.
- Content Generation: News summaries, code assistants.

## 6. Sample Code Snippet (Hugging Face Transformers)

from transformers import pipeline

# Zero-shot sentiment classification

classifier = pipeline("zero-shot-classification", model="facebook/bart-large-mnli")

text = "I love my new smartphone, battery life is amazing!"

labels = ["positive", "negative", "neutral"]

result = classifier(text, candidate\_labels=labels)

print(result)

# {'labels': ['positive', 'neutral', 'negative'], 'scores': [...]}

## 7. Common Interview Questions & How to Answer

• Q: "What's the difference between stemming and lemmatization?"

A: "Stemming crudely chops endings (e.g., 'running'→'run'), while lemmatization uses vocabulary and morphology to

return valid lemmas."

Q: "Why use subword tokenization?"

A: "Balances vocabulary size and out-of-vocabulary handling by representing rare words as subword units."

• Q: "How does self-attention work?"

A: "Each token computes attention scores via query/key dot-products against all tokens, then aggregates value vectors—capturing global context."

• Q: "How do you evaluate an NLP model?"

A: "Use task-specific metrics: accuracy/F1 for classification, BLEU/ROUGE for generation, exact match for QA."

• Q: "How do you mitigate bias in NLP?"

A: "Curate balanced corpora, apply debiasing techniques on embeddings, and audit outputs with fairness metrics."

### • Q: "What is Topic Modeling and how to optimize text for it?"

A: "Topic Modeling is an unsupervised machine learning technique used to automatically discover hidden topics in a collection of documents. The goal is to identify clusters of words (topics) that frequently occur together and assign them to documents in varying proportions.

To optimize text for modeling, next needs to be done:

Lowercasing: "The Team" → "the team"

Remove punctuation/special characters

Remove numbers

Remove stopwords: "the", "is", "and", etc.

Tokenization: Split into words

Lemmatization/Stemming: "running" → "run"

Remove dates

Remove websites

### Bag of Words (BoW)

Bag of Words (BoW) is a method of converting text into numerical features for use in machine learning algorithms.

Key Idea:

Treat a text document as a "bag" of individual words, ignoring grammar and word order, but keeping word
counts.

How it Works:

- 1. Build a vocabulary of all unique words across your dataset.
- 2. For each document:

- Count how many times each word from the vocabulary appears.
- o Represent the document as a **vector of word frequencies**.

Let's say we have 2 sentences:

Sentence 1: "I love NLP"

Sentence 2: "I love machine learning"

### Step 1: Vocabulary

All unique words = [I, love, NLP, machine, learning]

 $\rightarrow$  Vocabulary size = 5

## Step 2: Vectorize

Sentence	1	love	NLP	machine	learning
I love NLP	1	1	1	0	0
I love machine learning	1	1	0	1	1

Each sentence becomes a feature vector of word counts.

# Pros

- Simple, fast, easy to implement
- Works surprisingly well for many baseline NLP tasks
- Good with algorithms like Naive Bayes, Logistic Regression

# X Cons

- High dimensionality (huge vocabulary → sparse vectors)
- No semantics (e.g., "good" and "great" are treated differently)
- No word order (can't capture meaning based on structure)

## **Variants and Improvements**

TF-IDF (Term Frequency - Inverse Document Frequency)

Weights frequent words down and rare, informative words up:

$$ext{TF-IDF}(t,d) = ext{TF}(t,d) imes \log \left(rac{N}{ ext{DF}(t)}
ight)$$

#### Where:

- TF(t,d): term frequency in document
- DF(t): number of documents containing term
- N: total number of documents

#### n-grams

Use combinations of words:

- Unigram: ["I", "love", "NLP"]
- Bigram: ["I love", "love NLP"]
- Helps retain some word order

## Word Embedding and Word2Vec

1. Definition & Core Concept

Word Embeddings are dense vector representations of words in a continuous space, where semantically similar words lie close together. Word2Vec is a family of shallow, two-layer neural models that learn these embeddings by predicting word co-occurrence in a corpus.

- Skip-Gram: Given a target word, predict its context words.
- CBOW (Continuous Bag-of-Words): Given surrounding context words, predict the target word.

### 2. Why It Matters in NLP Pipelines

- Captures Semantics: Embeddings encode analogies ("king man + woman ≈ queen") and word similarities.
- Dimensionality Reduction: Transforms large sparse one-hot vectors into low-dimensional dense vectors (e.g., 100–300 D).
- Transfer Learning: Pre-trained embeddings can be plugged into downstream models—boosting performance with less data.

## 3. How Word2Vec Works

Component	Skip-Gram	CBOW
Objective	Maximize probability of context words given the center word.	Maximize probability of center word given context words.
Training	For each (center, context) pair, update embeddings via gradient descent on softmax (or negative sampling).	Similar but flips input/output roles.
Optimization	Often uses Negative Sampling or Hierarchical Softmax to approximate the full softmax for efficiency.	Same sampling tricks apply.

# 4. Advantages & Limitations

Advantages	Limitations
Efficient to train on large corpora	Static embeddings—no context sensitivity
Captures linear semantic relationships	Struggles with polysemy (one embedding per word)
Widely used as building block for many NLP tasks	Doesn't encode subword or character information

# 5. Real-World Applications & Examples

- Initialization: Pre-trained Word2Vec vectors initialize models for text classification or sequence tagging.
- Similarity Search: Find synonyms or related terms (e.g., in recommendation engines).
- Clustering & Visualization: Group related words or explore semantic spaces (e.g., via t-SNE).
- Analogy Tasks: Evaluate embedding quality with "man:king :: woman:?" quizzes.

# 6. Sample Code Snippet (Gensim)

from gensim.models import Word2Vec

from gensim.test.utils	import	common_texts
------------------------	--------	--------------

# Train a Skip-Gram model

model = Word2Vec(sentences=common\_texts, vector\_size=100, window=5, sg=1, negative=5, epochs=50)

# Access embeddings

vec\_king = model.wv['king']

# Find most similar words

similar = model.wv.most\_similar('king', topn=5)

print(similar)

### 7. Common Interview Questions & How to Answer

• Q: "Why use negative sampling instead of full softmax?"

A: "Full softmax requires costly normalization over the entire vocabulary; negative sampling updates only a few 'negative' examples per step—much faster."

• Q: "What's the difference between Word2Vec and GloVe?"

A: "GloVe is a count-based model that factorizes a global co-occurrence matrix, while Word2Vec is predictive, training a neural net to predict context."

• Q: "How do embeddings handle out-of-vocabulary words?"

A: "In Word2Vec you can't, but subword models like FastText break words into n-grams to form embeddings for rare or new words."

• 0: "Why might Word2Vec embeddings capture analogies?"

A: "Because the training objective encourages linear relationships among word vectors that mirror semantic relationships in text."

• Q: "When would you prefer contextual embeddings over Word2Vec?"

A: "If your task needs to disambiguate word senses or capture sentence-level nuance—then use BERT, ELMo, or Transformer-based models."

### Sequence-to-Sequence (seq2seq) Encoder-Decoder Neural Networks

1. Definition & Core Concept

Seq2Seq models map an input sequence (e.g., a sentence) to an output sequence (e.g., a translation), even if the two sequences differ in length.

They consist of two main components:

- Encoder: Compresses the input into a fixed-length context vector.
- Decoder: Generates the output sequence token by token from this context.

**Example**:

Input (English): "I love cats"

Output (French): "J'aime les chats"

## 2. Architecture Overview

 $Input \ Sequence \rightarrow [Encoder \ RNN] \rightarrow Context \ Vector \rightarrow [Decoder \ RNN] \rightarrow Output \ Sequence$ 

- Encoder: Processes input tokens into hidden states, compressing them into a fixed-size context vector (last hidden state).
  - Decoder: Takes this vector as input and predicts the next word in the output sequence iteratively.

# 3. Key Components

Component	Role
Encoder RNN	Reads the input sequence, updates hidden states.
Context Vector	Summarizes the entire input; passed to decoder as starting state.
Decoder RNN	Generates target sequence step by step using context + previous outputs.
Teacher Forcing	During training, the decoder receives ground-truth previous token, instead of its own prediction.

4. Enhancements Over Basic Seq2Seq

Attention Mechanism

- Rather than relying only on a single fixed context vector, attention allows the decoder to focus on different encoder hidden states at each time step.
- The attention mechanism helps the decoder focus on relevant parts of the input sequence when generating the output, particularly important for longer sequences.
  - Improves performance on long sequences.

# Bidirectional Encoder

• Processes input both forward and backward to capture richer context.

## Beam Search (Decoding)

• Keeps track of top k most likely sequences rather than just the most likely at each step.

### 5. Training and Loss

- Loss Function: Typically categorical cross-entropy between predicted tokens and ground-truth tokens.
- Optimization: Trained end-to-end using backpropagation through time (BPTT).

# 6. Applications

Task	Example
Machine Translation	English → Serbian, etc.
Text Summarization	Long article → Short summary
Speech Recognition	Audio waveform → Text
Question Answering	Context + question → Answer
Chatbots / Dialogue	User message → Bot response

# 7. Sample Code Snippet (Simplified - PyTorch)

```
# Encoder
class EncoderRNN(nn.Module):
  def __init__(self, vocab_size, hidden_size):
    super().__init__()
    self.embedding = nn.Embedding(vocab_size, hidden_size)
    self.rnn = nn.GRU(hidden_size, hidden_size)
  def forward(self, input):
    embedded = self.embedding(input)
    output, hidden = self.rnn(embedded)
    return hidden
# Decoder
class DecoderRNN(nn.Module):
  def __init__(self, hidden_size, output_size):
    super().__init__()
    self.embedding = nn.Embedding(output_size, hidden_size)
    self.rnn = nn.GRU(hidden_size, hidden_size)
    self.out = nn.Linear(hidden_size, output_size)
  def forward(self, input, hidden):
    embedded = self.embedding(input).unsqueeze(0)
    output, hidden = self.rnn(embedded, hidden)
    output = self.out(output.squeeze(0))
    return output, hidden
```

## 8. Interview-Worthy Questions & Answers

• Q: "Why does vanilla seq2seq struggle with long sequences?"

A: "Because the entire input is compressed into a single context vector—information bottleneck. Attention solves this."

• Q: "What is teacher forcing and when is it used?"

A: "During training, we feed the correct previous token into the decoder instead of its last prediction. It speeds convergence but can cause exposure bias."

• Q: "What's the difference between greedy decoding and beam search?"

A: "Greedy picks the most probable token at each step; beam search keeps multiple candidate sequences and explores combinations."

• Q: "How is attention computed in encoder-decoder models?"

A: "Each decoder step computes a score for every encoder hidden state, applies softmax to get weights, and forms a weighted sum for context."

### 9. Seq2Seq vs. Transformers

Seq2Seq RNN	Transformer
Sequential (can't parallelize)	Fully parallelizable via self-attention
Struggles with long range deps	Captures global dependencies effectively
Simpler, easier to interpret	Heavier but more powerful

### Attention

## 1. What Is Attention in Neural Networks?

Attention is a technique that allows models—especially in sequence tasks—to focus on the **most relevant parts of the input** when making predictions. Instead of compressing the entire input into one fixed-size vector (as in basic seg2seg), attention computes a weighted combination of all encoder states for each decoder step.

Think of it like a spotlight: the model learns where to look in the input for each output token.

### 2. Motivation

- Problem: Vanilla seq2seq models struggle with long inputs due to fixed-size bottlenecks.
- Solution: Attention dynamically decides which input tokens matter most during decoding.
- 3. New How It Works The Mechanics

At each decoder time step:

- 1. Query: The current decoder hidden state.
- 2. Keys & Values: All encoder hidden states.
- 3. Score (Alignment): Each encoder state is scored based on similarity to the query.
- 4. Weights: Softmax is applied to scores → attention weights.
- 5. Context Vector: A weighted sum of encoder states (values) based on attention weights.

[Decoder Hidden State]  $\rightarrow$  compare with  $\rightarrow$  [Encoder Hidden States]

1

Compute attention weights (via dot-product or MLP)

1

Context Vector = weighted sum of encoder outputs

4. 🔢 Math Formula (Dot-Product Attention)

Let:

- Q = Query (decoder hidden state)
- K = Keys (encoder hidden states)
- V = Values (encoder hidden states)

Then:

$$\operatorname{Attention}(Q,K,V) = \operatorname{softmax}\left(rac{QK^T}{\sqrt{d_k}}
ight)V$$

- dk = dimensionality of key vectors (used to scale scores)
- The softmax produces attention weights.

• The output is a weighted sum of the values V.

# 5. Types of Attention

Туре	Description
Additive (Bahdanau)	Uses a feedforward network to score relevance.
Dot-product (Luong)	Simpler, faster – uses the inner product of the query/key.
Self-Attention	Query, key, and value come from the same input sequence (used in Transformers).
Multi-Head Attention	Multiple attention mechanisms in parallel to capture different features.

# 6. <a> Applications</a>

- Machine Translation: Focus on relevant source words while decoding.
- Summarization: Attend key content of long documents.
- Transformers: The core operation in models like BERT, GPT, T5.
- Vision (ViT): Attend to image patches instead of convolution.

# 7. **Code Snippet (Simple Dot-Product Attention – PyTorch)**

import torch

import torch.nn.functional as F

def attention(query, key, value):

```
scores = torch.matmul(query, key.transpose(-2, -1)) / query.size(-1)**0.5
```

weights = F.softmax(scores, dim=-1)

context = torch.matmul(weights, value)

return context, weights

# 8. Interview Questions & Answers

• Q: Why do we use softmax in attention?

A: To turn raw alignment scores into a probability distribution over input positions.

• Q: Why is attention better than a fixed context vector?

A: It allows the decoder to dynamically access the most relevant encoder states, improving performance on long or complex inputs.

• Q: What's the difference between self-attention and encoder-decoder attention?

A: Self-attention relates a sequence to itself (e.g., in Transformers); encoder-decoder attention relates the input sequence to the output sequence.

Q: Why scale dot-product attention by dk\sqrt{d\_k}dk?

A: Without scaling, dot products can grow large and push softmax into regions with very small gradients, making training unstable.

## 9. 🧠 Final Thought

Attention isn't just a mechanism—it's the foundation of Transformer architectures and modern NLP advances like GPT, BERT, and T5. Understanding it well means you're truly fluent in cutting-edge AI.

### ELMo, BERT, Hugging Face and the others

ELMo (Embeddings from Language Models)

# ★ Summary:

- Developed by AllenNLP (2018)
- Produces contextualized word embeddings Based on bi-directional LSTM trained with a language modeling objective

### Key Concepts:

- Input: Word-level tokens
- Architecture: 2-layer BiLSTM trained on a large corpus
- Each word gets a vector that depends on its entire sentence context
- Different from static embeddings like Word2Vec or GloVe

### Output:

$$ext{ELMo}_i = \gamma \sum_{j=0}^L s_j h_{i,j}$$

- Combines all LSTM layers with learned weights
- hi,j: Hidden state of token i from layer j
- sj,γ: learned weights/scaling

### Key Aspects of ELMo:

## **Contextualized Embeddings:**

ELMo's primary advantage is its ability to capture the context of a word, allowing it to represent the same word with different vectors depending on its surrounding words.

### **Bidirectional LSTM:**

ELMo uses a bidirectional LSTM, which processes the input text from both forward and backward directions, allowing it to consider the entire sentence when generating embeddings.

### **Deep Learning:**

ELMo is a deep learning model, meaning it learns complex representations of words from data rather than being explicitly defined.

### Pre-trained on Large Corpora:

ELMo is pre-trained on a large text corpus, which enables it to capture a broad range of linguistic patterns and context.

### **Application in NLP:**

ELMo is widely used in various Natural Language Processing (NLP) tasks, including sentiment analysis, text classification, and question answering.

### How ELMo Works:

### 1. Input Embedding:

The input text is first converted into numerical representations (embeddings) using an embedding layer.

#### 2. Bidirectional LSTM Lavers:

The embeddings are then passed through a series of bidirectional LSTM layers, which process the input in both forward and backward directions.

### 3. Contextualized Embeddings:

The hidden states of the LSTM layers are used to generate contextualized embeddings, which capture the meaning of words in context.

### 4. Integration with downstream models:

These contextualized embeddings can then be used as features in other machine learning models for tasks such as text classification or question answering.

### **BERT (Bidirectional Encoder Representations from Transformers)**

### Summary:

- Developed by Google AI (2018)
- Transformer encoder-only architecture
- Learns deep bidirectional representations using masked language modeling (MLM)

It's designed to understand the context of words in text by considering both the preceding and following words, allowing it to better grasp the meaning of ambiguous language. BERT is trained on large amounts of unlabeled text data and is used in various NLP tasks, including sentiment analysis, question answering, and text generation.

### Key Features and Functionality:

### Bidirectional Context:

BERT's core strength is its ability to analyze text in a bidirectional manner, meaning it considers both the left and right context of each word.

### • Transformer Architecture:

BERT uses the Transformer architecture, a deep learning model that is highly effective at capturing complex relationships in text.

## • Pre-training and Fine-tuning:

BERT is pre-trained on massive amounts of data, and then fine-tuned for specific NLP tasks by adding task-specific layers.

### Versatility:

BERT can be applied to a wide range of NLP tasks, including text classification, sentiment analysis, machine translation, and more.

### How it works:

### 1. Pre-training:

BERT is trained on large datasets, such as Wikipedia and the BooksCorpus, using masked language modeling and next sentence prediction.

### 2. Masked Language Modeling:

BERT randomly masks some words in the input text and then tries to predict the original words based on the surrounding context.

## 3. Next Sentence Prediction:

BERT learns to predict whether one sentence follows another, which helps it understand the relationships between sentences.

### 4. Fine-tuning:

Once pre-trained, BERT can be fine-tuned on specific datasets for different NLP tasks, such as sentiment analysis or question answering.

### Applications:

- Sentiment Analysis: Analyzing the emotional tone of text.
- Question Answering: Finding answers to questions based on given text.
- Text Generation: Creating new text based on given prompts or inputs.
- Machine Translation: Translating text from one language to another.
- Search Engine Optimization: Improving search results by understanding the context of user queries.

# **Hugging Face**

# What is Hugging Face?

## Hugging Face is:

- A company building tools for NLP, ML, and generative Al
- The creator of the transformers Python library
- Home to 100,000+ pretrained models across NLP, vision, audio, and multimodal
- A platform for model sharing, training, inference, and collaboration

## transformers Library provides:

- Pretrained models (BERT, GPT, T5, RoBERTa, etc.)
- Tokenizers
- Pipelines for easy inference
- Support for PyTorch, TensorFlow, and JAX

## a datasets Library gives:

- Access to 10,000+ datasets (SQuAD, IMDB, MNLI, etc.)
- Efficient with memory and disk
- Compatible with PyTorch/TensorFlow

# Transformer Neural Networks, ChatGPT's foundation

## 1. What Is a Transformer?

A Transformer is a deep learning architecture introduced in the 2017 paper "Attention is All You Need". It replaces recurrence (used in RNNs and LSTMs) with self-attention mechanisms, enabling efficient parallel processing and powerful context modeling.

Transformers are the core architecture behind GPT, BERT, T5, etc.

### 2. Transformer Building Blocks

Architecture (Encoder-Decoder)

Originally designed for sequence-to-sequence tasks (e.g., translation), with two components:

Input Text  $\rightarrow$  [Encoder Stack]  $\rightarrow$  Context  $\rightarrow$  [Decoder Stack]  $\rightarrow$  Output Text

But GPT and ChatGPT use only the decoder stack in an autoregressive way.

# Nore Components

Component	Function
Self-Attention	Relates each token to every other token in the sequence.
Multi-Head Attention	Allows the model to focus on different relationships in parallel.
Positional Encoding	Adds sequence order information since there's no recurrence.
Feedforward Network	Applies transformation to each token position independently.
Layer Normalization	Stabilizes and speeds up training.
Residual Connections	Helps gradients flow better (skip connections).

# 3. Self-Attention Mechanism (Key Insight)

Each token creates:

- Query (Q)
- Key (K)
- Value (V)

Then, for every token, compute:

$$\operatorname{Attention}(Q,K,V) = \operatorname{softmax}\left(rac{QK^T}{\sqrt{d_k}}
ight)V$$

This allows each token to attend to other tokens based on similarity, capturing context.

# 4. 🔥 Multi-Head Attention

Instead of one attention calculation, multiple heads learn to focus on different parts of the sequence:

$$\operatorname{MultiHead}(Q,K,V) = \operatorname{Concat}(\operatorname{head}_1,\dots,\operatorname{head}_h)W^O$$

Where:

Each head is an independent attention mechanism:

$$head_i = Attention(QW_i^Q, KW_i^K, VW_i^V)$$

 $\boldsymbol{W}_{i}^{Q}, \boldsymbol{W}_{i}^{K}, \boldsymbol{W}_{i}^{V}$  : Learnable projection matrices for the i-th head

 $W^O$ : Output projection matrix after concatenation

## 5. Nositional Encoding

Since Transformers have no recurrence, positional encodings are added to input embeddings to give the model a sense of order.

Let:

- pos: position in the sequence (e.g., 0, 1, 2, ...)
- i: dimension index
- d: total embedding dimension (e.g., 512)

$$egin{aligned} ext{PE}_{(pos,2i)} &= \sin\left(rac{pos}{10000^{2i/d}}
ight) \ ext{PE}_{(pos,2i+1)} &= \cos\left(rac{pos}{10000^{2i/d}}
ight) \end{aligned}$$

Even dimensions use sine and odd dimensions use cosine

This adds unique patterns to token embeddings based on their position.

6. C GPT & ChatGPT: Decoder-Only Transformers

GPT-style models:

- Use only the decoder part of the transformer.
- Apply causal (masked) self-attention to ensure the model only sees previous tokens.

## ChatGPT:

- Fine-tuned version of GPT using Reinforcement Learning from Human Feedback (RLHF).
- Learns to generate coherent, safe, and helpful dialogue.

# 7. 🎇 Transformer vs. RNN

Feature	Transformer	RNN / LSTM
Parallelizable	✓ Yes	<b>X</b> No
Handles Long Contexts	✓ Better (attention)	X Poor (vanishing gradients)
Training Speed		Slow
Accuracy	Higher in NLP	Lower in modern benchmarks

# 8. @ Applications of Transformers

• Language Models: GPT-3, GPT-4, LLaMA

• Translation: T5, mBART

Summarization: BART

Question Answering: BERT, RoBERTa

• Chatbots: ChatGPT, Claude, Gemini

• Vision: Vision Transformers (ViT)

# 9. Common Interview Questions

• Q: Why are transformers better than RNNs for NLP?

A: They process sequences in parallel, use self-attention to model long-range dependencies, and scale better.

• Q: What is the role of positional encoding?

A: It injects sequence order into the model, since attention alone is permutation-invariant.

Q: Why do we need multi-head attention?

A: It allows the model to capture diverse relationships by attending to multiple parts of the sequence simultaneously.

Q: How is ChatGPT trained?

A: Pretraining on large text corpora  $\rightarrow$  supervised fine-tuning  $\rightarrow$  RLHF to align behavior with human preferences.

#### 10. Summary

Transformers are a type of neural network architecture that uses an attention mechanism to process sequential data like text or images, capturing relationships between elements and understanding context. They've become a cornerstone of natural language processing (NLP) and are also being applied in computer vision.

Key Features and Principles Summary:

#### Attention Mechanism:

Transformers rely on the attention mechanism, which allows them to weigh the importance of different input elements when making predictions. This is unlike traditional recurrent neural networks (RNNs) and convolutional neural networks (CNNs), which process data sequentially or locally.

### **Parallel Processing:**

The attention mechanism enables transformers to process elements of a sequence in parallel, making training faster and more efficient than RNNs.

Sequence-to-Sequence (Seq2Seq)

### Tasks:

Transformers are particularly well-suited for sequence-to-sequence tasks, where an input sequence is transformed into an output sequence, such as machine translation or text generation.

**Encoder-Decoder** 

### Structure:

Transformers often have an encoder-decoder architecture. The encoder processes the input sequence, and the decoder generates the output sequence.

### **Foundation Models:**

Transformers are considered foundation models, which are pre-trained on large datasets and can be fine-tuned for specific tasks.

# Decoder-Only Transformers, ChatGPTs specific Transformer



What Are Decoder-Only Transformers?

A decoder-only transformer is a streamlined version of the original Transformer architecture that uses only the decoder block to perform causal, autoregressive language modeling - predicting the next token in a sequence based on all previous ones.

This is the exact architecture behind GPT-2, GPT-3, GPT-4, and ChatGPT.

Decoder Block Architecture

Each decoder layer consists of:

- 1. Masked Multi-Head Self-Attention
- → Prevents "cheating" by blocking access to future tokens.
  - 2. Feedforward Neural Network (FFN)
- → Position-wise dense layers to transform token representations.
  - 3. Add & Norm (Residual Connections + LayerNorm)
- → Helps with gradient flow and stabilizes training.

## Layer Flow:

Input Embedding

+ Positional Encoding

→ Masked Multi-Head Self-Attention

→ Add & Norm

→ Feedforward Layer

→ Add & Norm

→ Output

Key Concept: Causal (Masked) Attention

Instead of full self-attention (like in encoders), decoder-only models apply causal masking so each token only attends to itself and past tokens:

$$Attention_{i,j} = 0$$
 if  $j > i$ 

This ensures left-to-right generation, essential for text prediction.

Training Objective: Next Token Prediction

The model learns a probability distribution over the vocabulary:

$$P(x_1, x_2, ..., x_T) = \prod_{t=1}^T P(x_t \mid x_1, ..., x_{t-1})$$

Training is done using cross-entropy loss between predicted and actual tokens.

### in How ChatGPT Works

ChatGPT is based on a decoder-only GPT model, fine-tuned with human feedback:

- 1. Pretraining: Predict next tokens on massive web text corpora.
- 2. Supervised Fine-tuning: Human-labeled prompts and completions.
- 3. RLHF (Reinforcement Learning from Human Feedback): Rank responses, train a reward model, then fine-tune with Proximal Policy Optimization (PPO).

This adds helpfulness, safety, and alignment to the raw language model.

# Positional Encoding

Since no recurrence exists, positional encodings are added to token embeddings to provide sequence order awareness:

- · Classic: sinusoidal functions.
- GPT-style: learnable position embeddings.

### Decoder-Only vs Full Transformer

Feature	Decoder-Only (GPT)	Full Transformer (BERT, T5)
Uses encoder?	<b>X</b> No	✓ Yes
Uses a decoder?	✓ Yes	(in encoder-decoder models)
Training objective	Autoregressive LM	Masked LM or Seq2Seq
Can access future tokens?	<b>X</b> No	✓ (depends on setup)
Use case	Generation	Classification, translation

#### Interview Highlights

Q: Why use only the decoder?

A: It's all you need for unidirectional text generation (next-token prediction). Encoders aren't necessary.

Q: Why mask future tokens?

A: To prevent the model from using future information during training — it must simulate true generative behavior.

Q: How is context handled in GPT?

A: Through self-attention over all previous tokens (up to a context window limit, e.g., 2k-128k tokens).

O: How is ChatGPT different from GPT?

A: ChatGPT is a fine-tuned, instruction-following version of GPT trained with RLHF to handle dialogue better.

## PyTorch-Style Pseudocode for Masked Attention

mask = torch.tril(torch.ones(seq\_len, seq\_len)) # Lower triangle

 $scores = (Q @ K.T) / sqrt(d_k)$ 

scores.masked\_fill\_(mask == 0, float('-inf')) # Mask future

weights = softmax(scores, dim=-1)

output = weights @ V

# Final Take

Decoder-only transformers are the powerhouse behind modern generative AI. They combine simplicity, scalability, and performance, especially when paired with massive datasets and alignment techniques like RLHF.

## Summary!!

Decoder-only Transformers are a type of neural network architecture that focuses on generating text sequences one token at a time, similar to how humans write or speak. They are particularly well-suited for tasks like language modeling, text completion, and creative writing, as they can predict the next word in a sequence based on the preceding words.

### Key Features:

#### Autoregressive

Decoder-only Transformers are designed to predict the next token in a sequence based on the tokens that have already been generated. This makes them ideal for tasks where the model needs to generate sequential data, like text.

### Masked Self-Attention:

Each layer in a Decoder-only Transformer uses a masked self-attention mechanism. This allows the model to attend to all previous tokens in the sequence when predicting the next token, but not the tokens that come after it.

#### Feed-forward Neural Networks (MLPs):

Each layer also includes a feed-forward neural network (or Multi-layer Perceptron, MLP). These MLPs process the output of the self-attention mechanism and refine the embeddings of each token.

### How it Works:

- Input: The model receives a sequence of input tokens.
- Token Embeddings: Each token is converted into a numerical representation called a token embedding.
- **Positional Encoding:** Positional encodings are added to the token embeddings to convey the order of the tokens in the sequence
- Stacked Decoder Layers: The input is passed through multiple stacked decoder layers.
- Self-Attention: Each layer uses self-attention to weigh the importance of different tokens when predicting the next token
- Feed-forward Networks: The output of the self-attention mechanism is then processed by a feed-forward network.

• Output: The final output is a probability distribution over the vocabulary, representing the likelihood of each token being the next in the sequence.

## **Encoder-Only Transformers (like BERT) for RAG**

What Is an Encoder-Only Transformer?

An Encoder-Only Transformer uses just the encoder stack of the original transformer architecture — meaning it processes input all at once, focusing on understanding and representing the input context.

Example: BERT (Bidirectional Encoder Representations from Transformers)

- Processes input bidirectionally (considers both left and right context).
- Trained with masked language modeling and next sentence prediction.
- Outputs rich contextual embeddings for each token (and the entire sentence via [CLS] token).

## Key Encoder-Only Characteristics

Feature	Encoder-Only (e.g., BERT)
Input directionality	Bidirectional
Output type	Contextual embeddings
Used for	Classification, retrieval, NER, etc.
Generation?	➤ No (not designed for generating text)

- Use Case: Retrieval-Augmented Generation (RAG)
- What is RAG?

RAG combines:

- 1. Retriever (Encoder-Only Transformer like BERT) → Finds relevant documents or passages.
- 2. Generator (Decoder-Only Transformer like GPT) → Uses retrieved info to generate the answer.

It enables language models to access external knowledge dynamically, reducing hallucinations and improving factual accuracy.

- Example: BERT + FAISS + GPT-3
  - 1. Encode your corpus using a BERT-like model (e.g., sentence-transformers).
  - 2. Index all document embeddings using FAISS (a fast similarity search library).
  - 3. At runtime, encode the query with BERT, find top-k similar passages.
  - 4. Feed those passages + query to a decoder-only model like GPT to generate the final answer.
- Why Encoder-Only for Retrieval?
  - Dense Vector Representations: BERT outputs embeddings that capture meaning, ideal for similarity search.
- Semantic Search: Unlike keyword search, it matches meaning ("Where can I buy sneakers?"  $\approx$  "Best shoe stores nearby").
- Dual Encoding: In some cases, both questions and documents are encoded (Bi-Encoder models) and compared using cosine similarity or dot product.

### ■ Encoder Variants for RAG

Model	Description	Use in RAG	
BERT	Bidirectional encoder; base choice	Semantic search	
SBERT	Sentence-transformers fine-tuned BERT	Faster, better embeddings	

DistilBERT	Lightweight BERT	More efficient retrieval
ColBERT	Late interaction for better accuracy	Large-scale retrieval

# ✓ Interview Cheat Sheet

• Q: Why can't BERT generate text?

A: It lacks a decoder and is not trained autoregressive — it's for understanding, not generation.

• Q: What makes BERT good for retrieval?

A: Its contextual embeddings capture semantic similarity, not just lexical overlap.

Q: What's the benefit of RAG?

A: It allows generative models to cite real facts instead of guessing, improving accuracy and trustworthiness.

• Q: How is information retrieved in RAG?

A: By encoding the query and corpus, then finding top-k similar passages using vector search.

### Code Sketch: Retrieval with BERT and FAISS

from sentence\_transformers import SentenceTransformer import faiss

model = SentenceTransformer('all-MiniLM-L6-v2')
corpus\_embeddings = model.encode(corpus, convert\_to\_tensor=True)

### # Build FAISS index

index = faiss.IndexFlatL2(corpus\_embeddings.shape[1])
index.add(corpus\_embeddings.cpu().numpy())

### # Query

query = "What is the capital of Serbia?"
query\_embedding = model.encode([query])
D, I = index.search(query\_embedding, k=5)
top\_docs = [corpus[i] for i in I[0]]

# **Missing Data Handling**

Handling missing data is crucial to ensure that machine learning models train effectively and avoid bias or errors.

# Types of Missingness:

- MCAR (Missing Completely At Random): No relationship between missing data and any variable.
- MAR (Missing At Random): Missingness depends on observed data.

• MNAR (Missing Not At Random): Missingness depends on unobserved data.

# **★**Simple Imputation Techniques

Technique	Description	Suitable For
Mean/Median/Mode	Replace missing values with column mean/median/mode	Numerical (Mean/Median), Categorical (Mode)
Constant Value	Fill missing values with a fixed number (e.g., -999)	Categorical/Numerical
Forward Fill / Backward Fill	Use previous/next value to fill missing data	Time Series Data

# Python example code:

from sklearn.impute import SimpleImputer

import numpy as np

X = np.array([[1, 2], [np.nan, 3], [7, 6]])

imputer = SimpleImputer(strategy='mean') # or 'median', 'most\_frequent'

X\_imputed = imputer.fit\_transform(X)

# **★**Advanced Imputation Techniques

Technique	Description	Tool
KNN Imputer	Fill missing values using the mean of k-nearest neighbors	sklearn.impute.KNNImputer

Iterative Imputer	Multivariate imputation — predicts missing value as a regression problem	sklearn.impute.IterativeImpute r
MICE (Multiple Imputation by Chained Equations)	Bayesian approach to iteratively impute missing data	fancyimpute, statsmodels
Deep Learning Autoencoders	Learn representation and reconstruct missing parts	TensorFlow / PyTorch

## Python example code:

from sklearn.impute import KNNImputer

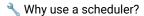
knn\_imputer = KNNImputer(n\_neighbors=3)

X\_knn = knn\_imputer.fit\_transform(X)

- Always analyze why data is missing.
- Impute only on training data, then apply the same transform to test data.
- Use pipelines to avoid data leakage during preprocessing.

## **Learning Rate Schedulers**

The **learning rate** controls how big each update step is during gradient descent. A **scheduler** adjusts this value over time to improve convergence and stability.



- Prevent overshooting the minimum early in training.
- Speed up convergence.
- Fine-tune learning in later stages.
- Escape local minima.

### \* Popular Options:

## ReduceLROnPlateau:

Monitors a validation metric (e.g., loss/accuracy). Reduces the learning rate when a metric has stopped improving.

from tensorflow.keras.callbacks import ReduceLROnPlateau

reduce\_lr = ReduceLROnPlateau(monitor='val\_loss', factor=0.5, patience=3, min\_lr=1e-6)

Parameter	Description
factor	Multiplicative factor of LR reduction
patience	Number of epochs with no improvement
min_lr	Lower bound for LR

## **Cosine Annealing:**

Learning rate decreases following a **cosine curve** from initial LR to a minimum. Good for cyclic training (warm restarts).

$$\eta_t = \eta_{min} + rac{1}{2}(\eta_{max} - \eta_{min})(1 + \cos(rac{T_{cur}}{T_{max}}\pi))$$

- T\_max: total number of iterations/epochs before restart.
- Often used with SGDR (Stochastic Gradient Descent with Warm Restarts).

from torch.optim.lr\_scheduler import CosineAnnealingLR

scheduler = CosineAnnealingLR(optimizer, T\_max=10, eta\_min=1e-6)

### **Exponential Decay:**

$$\eta_t = \eta_0 \cdot \exp(-k \cdot t)$$

• Learning rate decreases exponentially with time.

# **Step Decay:**

• Reduce LR by a factor every fixed number of epochs.

When to use what?

Scenario	Recommended Scheduler

Plateau in validation loss	ReduceLROnPlateau
Smooth decay needed	CosineAnnealing
Cyclical training or warm restarts	CosineAnnealing with restarts
Fast initial training, slow finish	Exponential or Step Decay

## **Gradient Clipping**

**Gradient Clipping** is a technique to **limit (clip)** the magnitude of gradients during backpropagation to prevent **exploding gradients**.

This is especially useful in:

- Recurrent Neural Networks (RNNs)
- Very deep neural networks
- Situations where gradients become unstable and training fails

## **How Clipping Works**

Two common strategies:

# Clip by value

• Clamp gradients element-wise to a range:

torch.nn.utils.clip\_grad\_value\_(model.parameters(), clip\_value=1.0)

# Clip by norm

• If the global norm of gradients exceeds a threshold, scale all gradients down proportionally:

torch.nn.utils.clip\_grad\_norm\_(model.parameters(), max\_norm=1.0)

## Norm clipping formula:

$$\mathbf{g}_{clipped} = \mathbf{g} \cdot \frac{\theta}{\|\mathbf{g}\|_2}$$

Where:

- g = gradient vector
- $\theta$  = threshold

When to use it:

- If you observe **exploding losses** or NaNs.
- When training LSTMs, GRUs, or Transformers.
- With large learning rates that might push weights too far.

# ★ In tensorflow (Keras):

optimizer = tf.keras.optimizers.Adam(clipnorm=1.0)

### **AdamW Optimizer**

**AdamW** is a variant of the **Adam optimizer** that **decouples weight decay (L2 regularization)** from the gradient update. This improves training stability and generalization.

AdamW stands for:

- Adam: Adaptive Moment Estimation (combines momentum + RMSProp)
- W: Weight decay done right (not added to gradients)

### Problem with Adam:

In traditional Adam, **L2 regularization** was added directly to gradients:

$$heta = heta - \eta \cdot (
abla_{ heta} L + \lambda \cdot heta)$$

This entangles optimization and regularization, which can lead to:

- Over-regularization
- · Reduced performance

## AdamW Solution:

AdamW decouples weight decay:

$$heta = heta - \eta \cdot 
abla_{ heta} L - \eta \cdot \lambda \cdot heta$$

Two separate steps:

- 1. Gradient update: uses only the loss gradient.
- 2. Weight decay: applies a small penalty directly to the weights.

## **MLflow**

MLflow is an open-source platform to track experiments, package models, and manage the full ML lifecycle.

## What does MLflow offer?

Component	Purpose	
Tracking	Log, compare, and visualize parameters, metrics, artifacts, and models.	
Projects	Define reproducible ML code (using conda/env/docker).	
Models	Deploy and serve models across platforms.	
Model Registry	Manage and version ML models centrally.	

# **General Usage:**

- Keeps track of model experiments and results.
- Helps manage hyperparameter tuning (learning rate, batch size, etc).
- Supports multiple backends (local, S3, databases).
- Integrates with frameworks like TensorFlow, PyTorch, XGBoost, LightGBM.

### Basic Workflow:

```
import mlflow
with mlflow.start_run():
mlflow.log_param("lr", 0.01)
mlflow.log_metric("accuracy", 0.92)
mlflow.log_artifact("model.pkl")
```

Then view results with: mlflow ui

### **Summary**

MLflow is used to manage the machine learning lifecycle from initial model development through deployment and beyond to sunsetting. It combines: Tracking ML experiments to record and compare model parameters, evaluate performance, and manage artifacts (MLflow Tracking)

### **Diffusion Models**

(Denoising Diffusion Probabilistic Models - DDPM)

Diffusion models are a class of generative models that learn to **generate data** (like images or audio) by reversing a gradual noising process.

### 📸 Idea

- Forward Process (Diffusion): Gradually add noise to the data over multiple steps until it becomes pure noise.
- Reverse Process (Denoising): Train a neural network to learn how to remove the noise step by step, recovering the original data.

Math behind it:

### Forward (Noising) Process:

- Given input data x<sub>0</sub> (e.g., image):
- At time step t, add Gaussian noise:

$$q(x_t|x_{t-1}) = \mathcal{N}(x_t; \sqrt{1 - \beta_t}x_{t-1}, \beta_t I)$$

Where:

- β\_t is the noise variance schedule (small at start, larger later).
- After many steps,  $x_T$  is nearly standard normal  $\mathcal{N}(0, \mathbb{I})$ .

### Reverse (Denoising) Process:

Train a model  $\varepsilon_{-}\theta(x_{-}t, t)$  to predict the noise:

$$p_{\theta}(x_{t-1}|x_t) = \mathcal{N}(x_{t-1}; \mu_{\theta}(x_t, t), \Sigma_{\theta}(x_t, t))$$

We optimize using a simplified loss:

$$L = \mathbb{E}_{x_0,\epsilon,t}[||\epsilon - \epsilon_{\theta}(x_t,t)||^2]$$

• This is a simple **MSE loss** between the actual noise and predicted noise.

### Intuition

Think of it like a video in reverse: you're learning to reconstruct an image from static noise, one frame at a time.

# How are diffusion models used?

- Image generation: DALL·E 2, Imagen, and Stable Diffusion use this.
- Inpainting: Fill in missing parts of an image.
- Super-resolution: Upscale low-resolution images.
- Audio: Used in voice synthesis (e.g., WaveGrad).

# Summary

Feature	Diffusion Models
Training Objective	Noise prediction (MSE loss)
Input	Noisy images
Output	Clean images
Architecture	U-Net with attention + time embedding
Best Use Case	Image/audio generation & enhancement